

MASTERARBEIT | MASTER'S THESIS

Titel | Title

Bidirectional Mass Spectrometry Modeling Aided by Auxiliary
Intermediate Steps From Graph Transformation Tools

verfasst von | submitted by
Philipp Reiser BSc

angestrebter akademischer Grad | in partial fulfilment of the requirements for the degree of
Master of Science (MSc)

Wien | Vienna, 2026

Studienkennzahl lt. Studienblatt | Degree
programme code as it appears on the
student record sheet:

UA 066 910

Studienrichtung lt. Studienblatt | Degree
programme as it appears on the student
record sheet:

Masterstudium Computational Science

Betreut von | Supervisor:

ao. Univ.-Prof. Mag. Dr. Christoph Flamm

Acknowledgements

I would like to thank my supervisor Christoph Flamm for the opportunity and his patience. For tolerating my occasional slow grasp of concepts and for all the support I received, I am grateful.

Gratitude is also extended to the entire TBI working group for their support and for creating a stimulating environment. Special thanks to Manuel Uhlir for getting me started and for the many helpful discussions and suggestions along the way.

Heartfelt thanks are due to my mother and grandmother, whose constant faith in me has carried me through many challenges. I also wish to acknowledge my grandfather for the memorable time I was able to share with him.

Finally, thank you, Aila, for your warmth and kindness at every step of this journey.

Contents

Abstract	vii
Zusammenfassung	viii
List of Tables	ix
List of Figures	xi
1 Introduction	1
1.1 Motivation	1
1.2 Objectives	2
1.3 Scope and Limitations	3
2 Background and Related Work	5
2.1 Fundamentals and Capabilities of Mass Spectrometry	5
2.2 Principles of Mass Spectrometry	7
2.2.1 Ionization and Fragmentation	8
2.2.2 Mass Analyzers	12
2.2.3 Detection	14
2.2.4 Manual Interpretation of the MS Spectrum	15
2.2.5 Reactions in Electron Ionization	16
2.3 Computational Mass Spectrometry	17
2.3.1 Library Search	18
2.3.2 In Silico Fragmentation	18
2.3.3 De Novo Structure Prediction	19
2.3.4 First Principles	19
2.3.5 Spectral Similarity and Distance Measures	20
2.4 Computational Representations of Molecules and Reactions	23
2.4.1 Molecular Graphs	23
2.4.2 Chemical Graph Transformation Systems	26
2.4.3 Graph Neural Networks for Molecular Representation Learning	29

3	Methodology	31
3.1	Datasets	31
3.1.1	Data Sources	31
3.1.2	Preprocessing - Generating the Synthetic Data	32
3.1.3	Data Splits	36
3.2	Latent Representation Schemes	37
3.2.1	Molecular Graph Encoding	38
3.2.2	Two-Level Molecular Latent Representation	38
3.2.3	Rule Representation and Hyperedge Annotation	40
3.2.4	Spectral Vector and Embedding Representation	40
3.2.5	Disadvantages	41
3.3	Model Architecture	41
3.3.1	Problem Formulation	43
3.3.2	Overview of the Learning Framework	44
3.3.3	Forward Model: Structure-to-Spectrum Prediction	47
3.3.4	Backward Model: Spectrum-to-Structure Retrieval	49
3.3.5	Bidirectionality	49
3.3.6	Loss Function	50
3.3.7	Training Procedure	52
4	Results and Discussion	53
4.1	Computation Times of Fragmentation Expansion	53
4.1.1	Visualization of Computation Times	53
4.2	Evaluation of Structure-Spectrum Translation	54
4.2.1	The Forward Model Is Molecule-Specific	55
4.2.2	Mass-Controlled Retrieval	55
4.2.3	Why Raw Retrieval Scores Are Not the Right Axis	56
4.3	Challenges	59
5	Future Work	60
5.1	From Fragment Encoding to Fragment Generation	60
5.2	Toward a Generate-Score Architecture	61
5.3	Advantages of Learning Fragment Generation	61
5.4	Scientific and Technical Challenges	63
5.5	A Plausible Development Path	63
	Bibliography	64
A	Appendix	71
A.1	Hyperparameter Tables and Configurations	71
A.2	Github Repository	73

A.3 AI Use Statement	73
--------------------------------	----

Abstract

Mass spectrometry is a cornerstone technique in analytical chemistry, metabolomics, pharmaceutical research, and many other fields. The prediction of mass spectra from molecular structures, and the inverse inference of molecular structures from spectra, remain active areas of research. While machine learning methods increasingly improve forward spectrum prediction from known structures, the inverse problem — predicting plausible molecular structures from a given spectrum — remains difficult, ambiguous, and poorly interpretable.

This work investigates the integration of auxiliary intermediate fragments generated by graph-transformation tools into machine learning pipelines for bidirectional translation between molecular structures and mass spectra. These graph-based intermediate representations provide an explicit fragmentation layer between molecules and spectra, intended to make the prediction process more transparent than purely black-box approaches.

The proposed framework systematically incorporates graph-based fragmentation intermediates into structure $\leftrightarrow$ spectrum translation models. By combining molecular representations, spectral information, and fragment-level graph structures, it establishes a functional bidirectional pipeline with an explicit intermediate fragmentation layer. In a proof-of-concept evaluation, the forward model predicts molecule-specific spectra whose peaks fall at physically realizable fragment masses by construction, and spectrum-to-structure retrieval is assessed under a mass-controlled protocol that treats precursor mass as the baseline to beat. Within a fixed mass window, learned retrieval exceeds the analytic random-chance baseline on held-out molecules — structural signal that a raw retrieval score would otherwise attribute to mass filtering alone. The principal contribution is the framework itself, together with a mass-controlled evaluation that separates learned structural signal from mere mass consistency, providing a concrete basis for future work on learned fragment generation.

Zusammenfassung

Massenspektrometrie ist eine Schlüsseltechnik in der analytischen Chemie, der Metabolomik, der pharmazeutischen Forschung und vielen weiteren Bereichen. Die Vorhersage von Massenspektren aus Molekülstrukturen sowie die inverse Ableitung von Molekülstrukturen aus Spektren bleiben aktive Forschungsbereiche. Während Methoden des maschinellen Lernens die Vorhersage von Spektren aus bekannten Strukturen zunehmend verbessern, bleibt das inverse Problem — die Vorhersage plausibler Molekülstrukturen aus einem gegebenen Spektrum — schwierig, mehrdeutig und nur eingeschränkt interpretierbar.

Diese Arbeit untersucht die Integration zusätzlicher Zwischenfragmente, die durch graphbasierte Transformationstools erzeugt werden, in Machine-Learning-Pipelines für die bidirektionale Übersetzung zwischen Molekülstrukturen und Massenspektren. Diese graphbasierten Zwischenrepräsentationen bilden eine explizite Fragmentierungsschicht zwischen Molekülen und Spektren und sollen den Vorhersageprozess dadurch transparenter machen als rein Black-Box-basierte Ansätze.

Das vorgeschlagene Framework integriert graphbasierte Fragmentierungszwischenstufen systematisch in Modelle zur Struktur $\leftrightarrow$ Spektrum-Übersetzung. Durch die Kombination molekularer Repräsentationen, spektraler Information und fragmentbasierter Graphstrukturen entsteht eine funktionsfähige bidirektionale Pipeline mit einer expliziten Fragmentierungszwischenschicht. In einer Machbarkeitsstudie (Proof of Concept) sagt das Vorwärtsmodell molekülspezifische Spektren vorher, deren Peaks konstruktionsbedingt bei physikalisch realisierbaren Fragmentmassen liegen; die Spektrum-zu-Struktur-Wiederauffindung (Retrieval) wird unter einem massenkontrollierten Protokoll bewertet, das die Vorläufermasse als zu übertreffende Referenz behandelt. Innerhalb eines festen Massenfensters übertrifft das gelernte Retrieval die analytische Zufallsbasislinie auf ungesesehenen Molekülen — ein strukturelles Signal, das eine rohe Retrieval-Kennzahl andernfalls allein der Massenfilterung zuschreiben würde. Der wesentliche Beitrag besteht im Framework selbst sowie in einer massenkontrollierten Evaluation, die gelerntes strukturelles Signal von bloßer Massenkonsistenz trennt und damit eine konkrete Grundlage für zukünftige Arbeiten zur gelernten Fragmenterzeugung bildet.

List of Tables

2.1	Sample Isotope Patterns	16
3.1	Notation used in the model description.	44
3.2	Loss Terms	51
A.1	Training hyperparameters. Tuned entries are the best trial of the Optuna sweep.	72

List of Figures

2.1	Simplified block diagram of a mass spectrometer	7
2.2	Simplified schematic of MS measurement principle	7
2.3	Schematic of EI ionization	9
2.4	Schematic of ESI ionization	10
2.5	Schematic of APCI ionization	11
2.6	Schematic of MALDI ionization	12
2.7	Schematic of an electron multiplier detector	15
2.8	Cosine similarity geometry	20
2.9	Wasserstein distance illustration	21
2.10	Wasserstein Upsampling Illustration	22
2.11	Cosine similarity versus Wasserstein distance as spectral comparison measures . .	24
2.12	Abstraction of a molecule as a graph	25
2.13	Double pushout graph rewriting	26
2.14	Example derivation hypergraph of a fragmentation	28
2.15	Graph-convolutional molecular encoding workflow	30
3.1	Schematic of the synthetic data-generation workflow	32
3.2	Flow diagram of predicate filter application	34
3.3	Heteroatom check predicate	35
3.4	Saturation check predicate as constrained path traversal	35
3.5	Alkyl check predicate	36
3.6	Overview of the Two-Level Representation	37
3.7	Two-Level Molecular Latent Representation	40
3.8	Training dataflow	46
3.9	Inference and retrieval	47
4.1	Fragmentation Computation Times	54
4.2	Mass-controlled retrieval: learned vs. chance	56

4.3	Cosine similarity and MRR for ablation variants	57
4.4	Recall@k for ablation variants	58
5.1	Future extension toward a Generate – Score architecture	62

Chapter 1

Introduction

1.1 Motivation

Mass spectrometry (MS) is one of the most widely used analytical techniques for molecular characterization across chemistry, metabolomics, environmental sciences, and pharmaceutical research. [1, p. 1] Its high sensitivity, reproducibility, and throughput have enabled large-scale molecular profiling in both research and industrial settings [2, p. 1764]. Modern high-resolution and tandem mass spectrometers generate high-throughput spectral data that can be compared against large repositories of reference spectra such as the NIST Mass Spectral Library to support compound identification and profiling [3, pp. 7274-7275].

Despite these advances in data acquisition, the interpretation of MS data remains a significant bottleneck [4, p. 1]. Structural elucidation from spectra is still largely expert-driven, relying on chemical intuition and established fragmentation rules [4, p. 11]. While thousands of spectra can be recorded within a short time frame [4, p. 1], the corresponding structural annotation remains comparatively slow and manual [4, p. 11], with very few untargeted studies identifying more than 20% of detected compounds [4, p. 3]. This imbalance creates a fundamental asymmetry between data generation and data analysis.

The intrinsic complexity of molecular fragmentation contributes substantially to this challenge [1, pp. 30-31]. Fragmentation processes are governed by chemically valid bond cleavages and rearrangements [1, pp. 30-31], leading to multiple competing fragmentation pathways. A single molecule may therefore generate numerous fragments through direct cleavages and rearrangement reactions [1, pp. 30-31]. Because fragmentation pathways depend sensitively on molecular structure and internal energy [1, p. 26, pp. 30-31], structurally distinct molecules may produce partially overlapping spectral signatures, while small structural variations can in-

duce significant spectral differences. [1, p. 240] Consequently, the inverse problem — inferring molecular structure from spectral data — can be regarded as inherently ill-posed.

Recent machine learning approaches aim to address this challenge by learning direct mappings between molecular representations and spectral outputs. Graph neural networks [5], sequence-based models [6, p. 966], and generative approaches have demonstrated promising performance in forward (molecule-to-spectrum) and inverse (spectrum-to-molecule) prediction tasks. However, purely data-driven approaches exhibit several limitations:

- They often lack mechanistic interpretability. [7, p. 3& 6]
- They do not explicitly model fragmentation pathways. [7, p. 3]
- They rely on statistical correlations rather than chemically structured priors. [6, p. 966]
- Their generalization to rare or out-of-distribution molecular scaffolds remains limited. [8, p. 12]

In particular, models that treat molecules as flat graphs or string encodings and spectra as unordered peak lists ignore the structured chemical processes underlying fragmentation. This omission reduces interpretability and may hinder robust extrapolation beyond the training distribution. [9, p. 3420]

Fragmentation in mass spectrometry is not an arbitrary stochastic process; it follows chemically constrained transformation sequences. [10, ch. 4& 8] Therefore, incorporating structured representations of fragmentation—such as derivation graphs encoding valid transformation rules—offers a principled way to introduce chemical inductive bias into machine learning models. By explicitly modeling fragment lineage and transformation pathways, it becomes possible to align molecular substructures with spectral evidence in a mechanistically meaningful manner.

This thesis is motivated by the hypothesis that integrating chemically structured derivation graphs into a machine learning framework can introduce chemical inductive bias and yield an explicit fragmentation layer, while supporting bidirectional structure-spectrum modeling. The aim is to bridge the gap between high-throughput spectral acquisition and scalable, chemically grounded interpretation by combining fragment-aware representations with data-driven learning.

1.2 Objectives

The objective of this thesis is to design, implement, and systematically evaluate a fragment-augmented machine learning framework for improving the bidirectional translation between

molecular structures and mass spectra. Specifically, the work investigates whether the integration of graph-based reaction and fragment derivation graphs, generated through chemically valid graph transformation rules, can serve as structured intermediate representations, and to what extent they affect predictive performance.

As mentioned in the introduction, elucidating structure from mass spectra is an inherently ill-posed inverse problem. This thesis therefore proposes the use of fragment-level intermediates derived from rule-based graph transformations to introduce a chemical basis to the learning process.

Rather than learning a direct mapping between molecular graphs and spectral representations, the proposed framework decomposes the task into a structured pipeline. In this pipeline, molecular graphs are first transformed into chemically plausible fragment sets. These fragment sets then mediate the relationship to spectral outputs. These fragment sets, together with molecular and spectral representations, are embedded into a shared latent space, enabling aligned embeddings across the domains.

Ultimately, this thesis assesses to what extent incorporating chemically valid graph-based fragment derivations into neural learning pipelines affects the predictive power of bidirectional structure-spectrum modeling in mass spectrometry, and outlines how such a framework could enable de novo molecular generation from spectral data.

1.3 Scope and Limitations

This thesis is restricted to electron ionization (EI) mass spectrometry, as EI fragmentation reactions are comparatively well understood and mechanistically characterized. [10, ch. 4& 8] This enables the application of chemically informed, rule-based graph transformations to generate fragment derivation graphs based on consistent assumptions. Other ionisation techniques are excluded due to their different fragmentation behaviour and the additional modelling complexity they would introduce.

A major limitation arises from the computational cost of graph-based fragment generation. Rule-based graph transformations lead to combinatorial growth in possible fragments, [11, p. 4] especially when multiple reaction rules and recursive fragmentation steps are considered. Ensuring chemical validity and avoiding duplicate graph states through isomorphism checking further increases computational demands. [11, p. 17]

Due to this scalability constraint, the experimental evaluation is limited to an order of magnitude of 100 compounds. While this is sufficient to demonstrate methodological feasibility, the restricted dataset limits chemical diversity and does not allow for generalisation. Consequently,

the results should be interpreted as proof of concept rather than as evidence of large-scale applicability.

It is also important to note that the reactions are defined in an ambiguous manner using graph symbols. Furthermore, the transformation rules do not cover all possible fragmentation pathways, and some may be incorrect. This means that the generated fragment derivation graphs may not capture all relevant chemical transformations, and the model’s performance may depend on the specific transformation rules chosen. The generated derivations are therefore constrained by the assumptions embedded in the rule set.

Chapter 2

Background and Related Work

2.1 Fundamentals and Capabilities of Mass Spectrometry

Mass spectrometry is a core analytical technique that characterizes chemical species by converting them into gas-phase ions and measuring their mass-to-charge ratio (m/z). [12, p. 5] A mass spectrometer therefore acts as a highly sensitive “molecular balance” whose output – a mass spectrum – reports which ionic species are present and with what relative abundances.

Conceptually, an MS experiment comprises three main steps:

- (i) **ionization**, in which neutral molecules are transferred into the gas phase and charged; [12, p. 1]
- (ii) **mass analysis**, in which ions are separated according to m/z ; and
- (iii) **detection**, in which ion signals are recorded and assembled into a spectrum. [12, p. 5]

Because many compounds form multiple ion species (e.g., adducts [12, p. 21], isotopologues [1, p. 67], fragments), the measured spectrum is not a direct image of the neutral molecule but a structured measurement that must be interpreted in the context of the ionization and fragmentation processes used. [12, p. 16-18]

A key capability of MS is the *identification of unknown compounds*. Even without prior knowledge of the analyte, MS provides constraints that narrow the space of possible identities. Accurate m/z measurements can restrict candidate molecular formulas; isotope patterns (e.g., from ^{13}C , ^{37}Cl , ^{81}Br) further refine these candidates by providing diagnostic abundance ratios and characteristic spacings. [1, p. 263] When fragmentation data are available – either through tandem MS (MS/MS) or through ionization methods that induce fragmentation – structural hypotheses can be tested by matching observed fragment ions and neutral losses to expected

substructures. [1, p. 225] In practice, unknown identification often combines these orthogonal cues with computational searches in spectral libraries and structure databases, where a measured spectrum is matched against reference spectra or predicted fragmentation patterns. [13, p. 8] For complex samples, chromatographic coupling (GC-MS or LC-MS) adds retention time as an additional dimension, [1, p. 475] improving selectivity and reducing confounding effects from co-eluting compounds. [12, p. 233-237]

Also in *purity control and quality assurance* MS serves a critical role, particularly in pharmaceutical manufacturing. By detecting the presence and abundance of impurities, degradation products, and unexpected byproducts, MS ensures that manufactured compounds meet specification and quality standards. The high selectivity and sensitivity of MS make it ideal for identifying trace contaminants that would escape less discriminating analytical techniques. Combined with chromatographic separation, MS provides definitive identification and quantification of both the target compound and potential impurities in a single analytical run. [14, p. 2]

MS is equally important for the *quantification of known substances*, where the goal is to estimate analyte concentration rather than molecular identity. [12, p. 260] Quantitative MS is grounded in the relationship between the number of detected ions and the amount of analyte introduced into the instrument. [12, p. 261 & 265-266] In idealized settings, ion signal intensity scales proportionally with concentration; in real measurements, this proportionality is skewed. Consequently, reliable quantification typically relies on calibration strategies. *External calibration* constructs a response curve from standards of known concentration, while *internal standardization* – often using isotopically labeled analogs – compensates for variability during sample preparation, injection, and ionization. [12, p. 264-268] The outcome is a method that can measure trace concentrations across wide dynamic ranges, which is central in many applications.

Beyond identity and amount, MS provides a systematic route to *deduce structural features of molecules*. Structural inference is primarily enabled by fragmentation behavior: when a precursor ion breaks apart, the resulting fragment ions reflect chemically plausible bond cleavages and rearrangements, thereby encoding information about functional groups, substructures, and connectivity. [15, p. i49] Certain fragmentation pathways are characteristic – for example, losses of small neutrals such as H₂O, NH₃, or CO can indicate the presence of hydroxyl, amino, or carbonyl functionalities, while specific fragment series can suggest alkyl chain lengths or aromatic substitution patterns. [1, p. 504] Importantly, structural elucidation from MS data is rarely straight forward; rather, it is an inference problem in which multiple candidate structures may explain the same spectral evidence, especially for isomers. [16, p. 53]

As a result, modern workflows increasingly integrate MS measurements with computational models that generate and rank candidate structures based on spectral consistency.

2.2 Principles of Mass Spectrometry




Figure 2.1: Simplified block diagram of a mass spectrometer. The sample is introduced into the ion source, where analyte molecules are ionized. The resulting ions are separated in the mass analyzer according to their mass-to-charge ratio (m/z), detected by the detector, and subsequently processed in the data analysis step. [17]




Figure 2.2: Simplified schematic of the measurement principle in a magnetic sector mass spectrometer. Following sample introduction, ionisation and acceleration generate an ion beam that is separated in the magnetic field by deflection that depends on m/z . The separated ions are then detected and transformed into an electrical signal for further processing. [18]

Instrument workflow

Figure 2.1 shows steps in a mass spectrometer: a *sample* is introduced, converted into gas-phase ions in the ion source, separated according to motion in the *mass analyzer* that depends on mass, detected in the *detector*, and finally interpreted during *data analysis*; Figure 2.2 depicts this sequence for a magnetic-sector instrument. The central observable is the mass to charge ratio m/z . A mass spectrum is therefore a distribution of signal intensity as a function of m/z , where the signal reflects the abundance of ions reaching the detector. [19, p. 1566] In most

practical contexts, intensities are reported as relative intensity, normalized such that the most intense peak equals 100% of the base peak. [19, p. 1550] The x-axis is conventionally labeled with m/z (mass to charge ratio), a dimensionless quantity. Labeling the axis with mass units such as *daltons* (Da), amu, or the unified atomic mass unit (u) is common in practice, although IUPAC recommends against it because m/z is a dimensionless quantity. [19, p. 1516]

Ion notation and precursor concepts

Ions observed in mass spectrometry are typically derived from a neutrally charged molecule M . In electron ionization (EI) a hard ionization method, removal of an electron yields the molecular ion [19, p. 1541] – also called the parent ion – denoted $[M]^{+\bullet}$, [19, p. 1582] a radical cation carrying both a positive charge and an unpaired electron. More generally, the term *precursor ion* refers to the ion selected for fragmentation, [19, p. 1578] from which product ions are generated. In soft ionization methods like electrospray, ion formation is often dominated by proton transfer, producing a protonated molecule $[M + H]^+$. [19, p. 1580]

Peak types: isotope, fragment, and base peaks

A measured spectrum contains several conceptually distinct peak types. *Isotope peaks* arise from naturally occurring heavier isotopes (e.g., ^{13}C , ^{37}Cl), producing characteristic $M + 1$, $M + 2$, ... patterns adjacent to the monoisotopic ion peak. *Fragment peaks* arise when precursor ions undergo dissociation, yielding smaller product ions whose m/z values and relative intensities encode structural information. The *base peak* is the most intense peak in the spectrum [19, p. 1524] and is used for normalization; [19, p. 1550] it is not necessarily the molecular ion and often corresponds to a particularly stable fragment. [20, p. 2] [1, pp. 258 & 263]

2.2.1 Ionization and Fragmentation

Ionization techniques are broadly classified into two categories based on the energy imparted to the analyte: *hard ionization* methods, which deliver high internal energy and promote extensive fragmentation, [19, p. 1547] and *soft ionization* methods, which minimize fragmentation and favor molecular ion preservation. [19, p. 1591] [21, p. 34]

Hard Ionization methods

Electron Ionization (EI) is a hard ionization technique in which a stream of high-energy electrons (typically 70 eV) collides with neutral molecules in the gas phase. [1, p. 197] The collision transfers energy to the molecule, causing the ejection of an electron and formation of a

radical cation $[M]^{\bullet+}$ with both a positive charge and an unpaired electron. [1, p. 194] Because the internal energy of the molecular ion typically exceeds the activation barriers for bond cleavage, the $[M]^{\bullet+}$ ion spontaneously fragments into smaller product ions and neutral molecules. [22, p. 3] As a result, EI mass spectra are characterized by rich fragmentation patterns that reflect the molecular structure. [23] A schematic of the EI source is shown in Figure 2.3.




Figure 2.3: Schematic illustration of the electron ionization (EI) process. Gaseous sample molecules (M) enter the source and are crossed by a beam of high-energy electrons (~ 70 eV) emitted from a heated filament and collected at the electron trap (anode). Collision with the electron beam ejects an electron from each molecule, forming a radical cation $[M]^{\bullet+}$. A repeller pushes the resulting ions toward the extraction and focusing lenses, which shape them into an ion beam directed to the mass analyzer. [24, p.18]

These characteristics make electron ionization particularly valuable for structural elucidation. [25] EI is the dominant ionization method in gas chromatography-mass spectrometry (GC-MS) and remains widely used in small-molecule analysis. [26, p. 1]

Spark Ionization (SI) , is a hard ionization technique, that often utilized to detect trace elements at parts-per-billion (ppb) levels. [27, p. 428] It uses high-voltage up to 100 kV pulsed electric sparks between two electrodes to vaporize, atomize, and ionize conductive or semi-conductive materials. [27, p. 428]

Soft Ionization methods

Electrospray Ionization (ESI) is a prominent soft ionization technique particularly well-suited for coupling with liquid chromatography (LC-MS) and for analysis of large biomolecules.

[28] In ESI, a liquid sample is converted into a fine aerosol of charged droplets by applying a high electric potential (typically several kilovolts) between a capillary needle and a counter electrode. [29, p. 3] As the droplets travel toward the mass analyzer, they undergo solvent evaporation and coulombic fission, eventually releasing gas-phase ions into the vacuum of the spectrometer. [29, p. 7] [24, p. 53]

ESI ionization is exceptionally gentle, with minimal or no fragmentation in the ion source. [29, p. 7] This characteristic makes ESI ideal for determining accurate molecular weights of intact molecules. [24, p. 50] Notably, ESI can generate multiply charged ions $[M + nH]^{n+}$ from large molecules, reducing their effective m/z and enabling their analysis on instruments with limited mass range. [24, p. 50]

A curtain gas is often employed to shield the ion path, preventing neutral contaminants and solvent droplets from reaching the mass analyzer while allowing ions to pass through. [12, p. 42] The sequence from charged droplet formation to gas-phase ion release is shown in Figure 2.4.




Figure 2.4: Schematic illustration of the electrospray ionization (ESI) process. A sample solution is dispersed by nebulizer gas and subjected to a high electric potential, producing a spray of charged droplets. Progressive solvent evaporation and droplet fission lead to the formation of gas-phase ions, which are transferred through the interface region into the mass detector. The figure highlights the transition from charged droplets to molecular ions, exemplified here by protonated species $[M + H]^+$.

[21, p.39]

Atmospheric Pressure Chemical Ionization (APCI) is a soft ionization method used primarily in LC-MS, [21, p. 39] where analytes are ionized at atmospheric pressure via ion-molecule chemistry initiated by a corona discharge. [19, p. 1523] In contrast to ESI, no high voltage is applied to the inlet capillary; instead, the mobile phase is nebulized, rapidly evaporated in a heated region, [24, p. 171] and the resulting vapor passes a high-voltage needle that generates reagent ions. Solvents dominate the early ionization steps [21, p. 39] and effectively act as a

reagent gas, transferring charge to neutral analytes mainly by proton transfer to form $[M + H]^+$.

Because APCI imparts relatively little internal energy, spectra typically show limited fragmentation and are well suited to quantitative LC-MS/MS, [24, p. 46] where signal is concentrated into a few precursor-ion peaks. A heated transfer tube and/or countercurrent curtain gas is used to decluster solvent-analyte adducts prior to ion extraction. [21, p. 39] APCI is often more efficient than ESI for less polar compounds, [1, p. 466] but its higher source temperatures can induce thermal degradation. [21, p. 39] A schematic of the APCI process is shown in Figure 2.5.




Figure 2.5: Schematic illustration of atmospheric pressure chemical ionization (APCI). The sample solution is nebulized and vaporized, and a corona discharge operated at high voltage generates reagent ions from the solvent. In the example shown, protonated ions transfer a proton to the analyte molecule M , forming the protonated molecular ion $[M + H]^+$, which is then introduced through the interface region into the mass detector.

[21, p.39]

MALDI (Matrix-Assisted Laser Desorption/Ionization) [30] is a soft ionization technique [31]. that uses a laser to desorb and ionize analytes from a solid matrix. [24, p. 36] It is particularly effective for large biomolecules and polymers, [12, p. 33][24, p. 36] producing predominantly singly charged ions. [12, p. 35]. Tuning laser energy and matrix composition can optimize ionization efficiency and control fragmentation. A general schematic of the MALDI process is shown in Figure 2.6.




Figure 2.6: Schematic illustration of matrix-assisted laser desorption/ionization (MALDI). A crystallized matrix containing embedded analyte molecules is irradiated with UV laser light at the sample stage, leading to desorption and ionization of matrix and analyte species. During plume formation, matrix ions and neutrals promote the formation of analyte ions, which are subsequently transferred into the mass analyzer.

[21, p. 42]

2.2.2 Mass Analyzers

After ionization, the generated ions must be separated according to their mass-to-charge ratio m/z before detection. This separation is performed by the *mass analyzer*. [24, pp. 5-6] Different analyzer designs exploit electric and magnetic fields to distinguish ions with different m/z values. Among the most widely used analyzers are magnetic sector instruments, time-of-flight (TOF) analyzers, and quadrupole analyzers. [24, pp. 108-109]

Magnetic Sector Mass Analyzers

Magnetic sector analyzers separate ions according to their motion in a magnetic field. The motion of a charged particle with mass m and charge q in electric and magnetic fields is described by the Lorentz force equation

$$\left(\frac{m}{q}\right) \mathbf{a} = \mathbf{E} + \mathbf{v} \times \mathbf{B},$$

where $\mathbf{E}$ is the electric field, $\mathbf{B}$ the magnetic field, $\mathbf{v}$ the ion velocity, and $\mathbf{a}$ the acceleration. Ions with identical mass-to-charge ratios experience identical forces and therefore follow identical trajectories inside the analyzer.

When ions enter a magnetic field perpendicular to their velocity, the Lorentz force bends their trajectory into a circular path. The balance between the magnetic force and the centripetal

force yields

$$qvB = \frac{mv^2}{r},$$

with r being the radius of the circular path, which leads to

$$\frac{m}{q} = \frac{B^2 r^2}{2V},$$

assuming the ions were accelerated through a potential V . Ions with different m/z therefore follow trajectories with different radii and can be spatially separated. [32, pp. 1799-1800]

Time-of-Flight (TOF) Analyzers

Time-of-flight analyzers separate ions according to the time required to travel a fixed distance after acceleration. When ions are accelerated through a potential V , they acquire kinetic energy

$$\frac{1}{2}mv^2 = qV.$$

Solving for the velocity gives

$$v = \sqrt{\frac{2qV}{m}}.$$

In a field-free drift region of length L , the flight time is therefore

$$t = \frac{L}{v} = L\sqrt{\frac{m}{2qV}}.$$

Consequently, lighter ions reach the detector earlier than heavier ions, allowing determination of their m/z values. [33, pp. 1150–1151]

Quadrupole Mass Filter (QMF)

Quadrupole analyzers consist of four parallel rods to which a combination of radio-frequency (RF) and direct current (DC) voltages is applied. [34, p. 2] The resulting oscillating electric field produces ion trajectories that are stable only for ions within a specific range of mass-to-charge ratios. Ions with stable trajectories pass through the quadrupole to the detector, while ions with unstable trajectories collide with the rods and are removed. [35, p. 1] By scanning the applied voltages, ions of increasing m/z can be transmitted sequentially, generating a mass spectrum. [34, p. 2]

Quadrupole Ion Trap (QIT)

The quadrupole ion trap operates on the same fundamental principle as the linear quadrupole analyzer described above, added with the ability of oscillating radiofrequency fields to selectively confine ions according to their mass-to-charge ratio. [36, p. 1156] The key distinction lies in the geometry of the electric field. The linear quadrupole uses four parallel rods to create a two-dimensional confining potential that allows ions to pass through axially. In contrast, the ion trap uses a three-dimensional quadrupolar field configuration to confine ions within a central cavity in all spatial dimensions. [37, p. 962] This fundamental geometric difference enables the QIT to accumulate and store ions over time [37, p. 973], rather than merely filtering them during transit, which has profound implications for sensitivity, scan modes, and the ability to perform multiple stages of mass analysis within a single device. [37, p. 984]

2.2.3 Detection

After ions have been separated by the mass analyzer, they must be detected and converted into a measurable electrical signal. [24, pp. 5-6] The detector is transforming the arrival of individual ions into quantifiable data that can be assembled into a mass spectrum. [24, p. 103]

The dominant detector in Electron Ionization Mass Spectrometry is the Electron Multiplier (EM). [24, pp. 103–104]

An electron multiplier operates by converting a single incoming ion into a large cascade of electrons through the following sequence:

1. An ion hits the first dynode
2. Secondary electrons are emitted
3. These electrons are accelerated to the next dynode
4. The process repeats (10-20 stages)
5. Result: a measurable current pulse

This cascade amplification principle is illustrated in Figure 2.7.




Figure 2.7: Schematic illustration of an electron multiplier (EM) detector. An incoming ion strikes the first dynode, causing the emission of secondary electrons. These electrons are then accelerated to subsequent dynodes, where they generate additional electrons in a cascading effect. After multiple stages, this results in a significant amplification of the original ion signal, producing a measurable current pulse that can be digitized and recorded as part of the mass spectrum.

[21, p. 50]

The amplification factor of an electron multiplier can reach 10^6 to 10^8 , enabling the detection of individual ions with high sensitivity. The resulting current pulses are digitized and recorded as a function of m/z , producing the final mass spectrum that serves as the basis for subsequent data analysis. [24, p. 104]

2.2.4 Manual Interpretation of the MS Spectrum

Data analysis transforms the raw spectrum into structural hypotheses: isotope patterns constrain elemental composition, [1, p. 67] while fragment ions and neutral losses constrain connectivity. [38]

The following summarize a typical manual workflow [24, pp. 215-216] for interpreting EI

mass spectra.

1. Determine the molecular ion peak (M) and its isotopic pattern
2. find isotope patterns (e.g., Cl, Br, S)

Table 2.1: Sample Isotope Patterns

Element	Diagnostic pattern
Cl	3:1 doublet (M, M+2)
Br	1:1 doublet
S	small M+2 peak
C	^{13}C gives predictable M+1

3. determine the base peak (100%) which is the energetically favored fragmentation
4. analyze fragment differences and identify neutral losses
5. match fragmentation to chemical logic defined by reaction rules

These practical heuristics are grounded in specific EI reaction mechanisms, which are outlined in the following subsection.

2.2.5 Reactions in Electron Ionization

Fragmentation reactions in electron ionization mass spectrometry follow well-characterized patterns that are not random but are reproducible pathways determined by [23] the internal energy distribution and the intrinsic chemistry of the molecular ion. [39, p. 846] Understanding these reaction mechanisms is essential for interpreting EI spectra and for predicting fragmentation patterns from molecular structure. [39, p. 846]

They have been studied extensively by McLafferty and Tureček, see [10]. The following subsections describe some of the most important mechanisms.

Homolytic single-bond cleavage

The most straightforward fragmentation pathway is *simple bond cleavage*, in which a single bond in the molecular ion breaks homolytically, yielding a charged fragment and a neutral radical. The probability of cleavage at a particular bond depends primarily on bond strength: weaker

bonds are more readily broken. [1, p. 30] For aliphatic hydrocarbons, C–C bonds adjacent to branching points or heteroatoms are particularly susceptible to cleavage. [1, p. 258]

Alpha Cleavage

Alpha cleavage is a characteristic fragmentation mode in molecules containing functional groups with heteroatoms. [24, pp. 220–221] In this process, the bond directly adjacent to the functional group is cleaved, producing a resonance-stabilized cation. [24, pp. 220–221] For instance, in ketones and aldehydes, cleavage of the C–C bond next to the carbonyl group generates an acylium ion $\text{R}-\overset{+}{\text{C}}=\text{O}$, which is stabilized by the resonance structure $\text{R}-\text{C}\equiv\overset{+}{\text{O}}$. [1, p. 230] This reaction pathway often produces intense peaks in the spectrum and serves as a diagnostic indicator for carbonyl-containing compounds. [1, pp. 234–235]

McLafferty Rearrangement

The *McLafferty rearrangement* is a hydrogen-transfer reaction characteristic of carbonyl compounds and other systems with abstractable γ -hydrogens.[24, p. 223] In the commonly depicted mechanism, a six-membered transition state facilitates the transfer of a hydrogen atom from the γ position to the carbonyl oxygen, accompanied by cleavage of the bond between the α and β carbons. [1, p. 265] The result is an enol radical cation and a neutral alkene. [24, p. 224] This rearrangement is particularly valuable for structural elucidation because it reveals connectivity information that would not be apparent from simple cleavage alone. [24, p. 224]

Charge-Induced Fragmentations

Beyond the radical-site-driven cleavages described above, fragmentation can also be initiated by the charge center itself. In such *charge-directed fragmentations*, the positive charge polarizes nearby bonds, facilitating heterolytic cleavage and producing even-electron cations and neutral radicals. [12, pp. 281]

2.3 Computational Mass Spectrometry

Computational tools significantly enhance MS capabilities for structure elucidation and compositional analysis. [13, pp. 1–2] Modern workflows integrate *in silico* fragmentation models, spectral database searches, and machine-learning methods to accelerate compound identification and enable high-throughput analysis. [23, p. 1] [13, p. 2]

In this section, we outline the key computational strategies in mass spectrometry, including library matching, *in silico* prediction, and *de novo* structure inference, before turning to their

mathematical and algorithmic foundations.

2.3.1 Library Search

Library search is a spectrum-matching procedure used to identify an unknown compound by comparing its measured spectrum against a reference database of known spectra. [13, p. 8] Each entry in such a database typically stores a reference spectrum together with metadata such as molecular formula, structure, and acquisition conditions.

The comparison operates on two observable quantities: the m/z values of detected peaks and their relative intensities. A query spectrum is considered a match if a reference spectrum reproduces both the positions and the abundance pattern of its peaks sufficiently well. [21, p. 35]

A typical library search workflow proceeds as follows: [13, pp. 1, 5]

- (i) **Pre-processing:** the query spectrum is converted into a form suitable for comparison, for example by peak selection, intensity scaling, or related preprocessing steps.
- (ii) **Candidate filtering:** database entries may be restricted using precursor mass or precursor m/z information within a tolerance window, reducing the number of candidates that must be evaluated. [13, p. 8]
- (iii) **Scoring and ranking:** each candidate reference spectrum is assigned a similarity score against the query, and candidates are ranked accordingly. [13, p. 5]

The most widely used similarity metric is the *cosine similarity*, computed over the vector of peak intensities aligned by m/z : [40, p. 1]

$$\cos(\mathbf{q}, \mathbf{r}) = \frac{\mathbf{q} \cdot \mathbf{r}}{\|\mathbf{q}\| \|\mathbf{r}\|}, \quad (2.1)$$

where $\mathbf{q}$ and $\mathbf{r}$ are the intensity vectors of the query and reference spectrum, respectively. A score of 1 indicates a perfect match; lower values indicate increasing dissimilarity.

2.3.2 In Silico Fragmentation

In silico fragmentation predicts how a molecule may dissociate in a mass spectrometer and generates a theoretical spectrum from these simulated fragment ions. [7, p. 1] The standard workflow starts from a candidate molecular structure, applies computational fragmentation rules or learned models to produce plausible fragments, converts these fragments into a predicted spectrum, and finally compares the prediction against the measured spectrum using a similarity score. [7, p. 2] This approach is particularly useful when no exact match is available in spectral libraries, because it enables structure-aware candidate ranking beyond pure reference lookup.

[13, pp. 1–2][23, p. 1] At the same time, prediction quality depends strongly on how well the model captures real instrument conditions, ionization-dependent chemistry, and fragment intensities; consequently, performance can degrade for compounds or fragmentation pathways that are underrepresented in the training data or incompletely described by rule sets. [41, p. 1989]

When these *in silico* predictions remain ambiguous, the workflow may shift to *de novo* inference directly from spectral constraints.

2.3.3 De Novo Structure Prediction

De novo structure prediction denotes the inference of molecular structure from scratch, i.e., without relying on pre-existing templates or premade reference repositories [42, p. 1]. Importantly, *de novo* characterizes the *template-free* nature of the workflow and does not, by itself, imply a specific modeling paradigm (e.g., physics-based versus data-driven). Accordingly, *de novo* approaches may incorporate empirical heuristics, statistical learning, or mechanistic constraints, depending on the application. [43, pp. 701–702]

2.3.4 First Principles

In contrast, *first-principles* modeling emphasizes deriving predictions without empirical parameterization, i.e., without fitting to experimental data. [44, p. 1] In physics and chemistry, such first-principles treatments are often realized via *ab initio* electronic-structure methods.

Ab Initio MS Prediction

Ab initio MS prediction refers to computational modeling of mass spectra directly from molecular structure using nonempirical quantum-mechanical methods that are, in principle, independent of experimental fitting [45, p. 1921]. In this sense, *ab initio* approaches constitute first-principles modeling, because they aim to predict fragmentation behavior and resulting spectra from the underlying physics and chemistry of the molecule rather than from learned correlations. Operationally, *ab initio* MS prediction typically relies on quantum-chemical calculations (to obtain energetics and feasible fragmentation channels) combined with dynamical or statistical treatments to relate accessible pathways to predicted fragment distributions.

While *ab initio* methods can provide mechanistic insight and may achieve high accuracy for small molecules, their computational cost often limits scalability to large molecular libraries and complex systems. [43, p. 702]

2.3.5 Spectral Similarity and Distance Measures

In computational mass spectrometry, evaluation metrics quantify how well predicted spectra match experimental observations. The choice of metric matters beyond diagnostics: when a metric is reused as a training objective, its geometry determines the loss landscape and therefore what a model is rewarded for reproducing (Section 3.3.6). Two families of comparison are available: *geometric* measures such as the cosine similarity, and *transport* measures such as the Wasserstein distance. The cosine similarity treats each spectrum as a vector and compares two such vectors A and B by the angle θ between them, normalising away their magnitudes so that only their relative direction enters the comparison; Figure 2.8 illustrates this geometry and eq. (2.2) gives the corresponding formula. [46, p. 4]




Figure 2.8: Two vectors A and B drawn from a common origin O . The cosine similarity is the cosine of the angle θ between them, $\cos \theta = (A \cdot B) / (\|A\| \|B\|)$, and equals the length of the projection of B onto the direction of A (amber) normalised by $\|B\|$. It depends only on the vectors' *directions*: scaling either vector changes its length but leaves θ , and hence the similarity, unchanged.

$$\text{CosSim}(A, B) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_i A_i B_i}{\|A\| \|B\|} \quad (2.2)$$

Because the comparison in eq. (2.2) is normalised by both vectors' norms, Cos_{Sim} depends only on direction and is invariant to the overall scale of either spectrum.

The Wasserstein distance, also known as the Earth Mover's Distance (EMD), measures the minimum cost of transporting one distribution onto the other, where cost is mass multiplied by the distance it is moved. [47, pp. 1–8] [48] For two equal-length vectors this reduces to sorting both and summing the absolute differences of matched ranks, a quantity that admits a simple geometric reading. Figure 2.9 draws each sorted vector as a cumulative staircase and shades the horizontal gap between them at every rank; the distance is the total shaded area, normalised by the number of elements.




Figure 2.9: Two equal-length vectors, sorted $\tilde{P} = [1, 6, 6, 10]$ (red) and $\tilde{Q} = [4, 4, 9, 12]$ (blue), are drawn as cumulative staircases: the horizontal axis is the value and the vertical axis is the rank index i/n , so each curve steps up by $1/n$ at every element. At each level the hatched amber rectangle spans the horizontal gap between the two staircases, whose width equals the matched-pair difference $|\tilde{P}_i - \tilde{Q}_i|$. The Wasserstein distance is the total hatched area divided by n (see eq. (2.3)), illustrating that the distance accumulates the cost of horizontally transporting one sorted distribution onto the other.

$$\text{WD}(P, Q) = \frac{1}{n} \sum_{i=1}^n |\tilde{P}_i - \tilde{Q}_i| = \frac{3 + 2 + 3 + 2}{4} = 2.5 \quad (2.3)$$

The matched-pair form in eq. (2.3) presupposes that both vectors contain the same number of elements, since each rank in $\tilde{P}$ must be paired with a corresponding rank in $\tilde{Q}$. Real spectra rarely satisfy this: two spectra typically differ in the number of detected peaks. To recover a matched-pair comparison without discarding or interpolating peaks, each vector is upsampled to a common length equal to the least common multiple of the two cardinalities, repeating every element enough times to preserve its relative share of the total mass. Figure 2.10 shows this construction for vectors of length four and three.




Figure 2.10: The inputs have different lengths, $P = [1, 3, 4, 5]$ ($n = 4$) and $Q = [1, 3, 5]$ ($m = 3$), so they are first brought to a common length before matched-pair differences can be formed. Each is upsampled to $\text{LCM}(4, 3) = 12$ by repeating every element of P three times and every element of Q four times (red and blue braces), giving aligned vectors $P' = [1, 1, 1, 3, 3, 3, 4, 4, 4, 5, 5, 5]$ and $Q' = [1, 1, 1, 1, 3, 3, 3, 3, 5, 5, 5, 5]$. The amber badges mark the four index positions where the upsampled entries differ, with per-index gaps 2, 1, 1, 1. The Wasserstein distance is the mean of these absolute differences (see eq. (2.4)), showing how repetition lets each value contribute mass in proportion to its share of the original vector, so distributions of different cardinality can be compared directly.

$$P' = [1, 1, 1, 3, 3, 3, 4, 4, 4, 5, 5, 5]$$

$$Q' = [1, 1, 1, 1, 3, 3, 3, 3, 5, 5, 5, 5]$$

$$\begin{aligned}
\text{WD} &= \frac{1}{12} \cdot (0 + 0 + 0 + 2 + 0 + 0 + 1 + 1 + 1 + 0 + 0 + 0) \\
&= \frac{5}{12} \\
&\approx 0.417
\end{aligned}
\tag{2.4}$$

The two measures differ in what they consider when two peaks fail to coincide. Because the cosine similarity accumulates intensity only at *shared* coordinates, it is blind to how far apart two peaks are: a peak displaced by a fraction of a Dalton and one displaced across the whole axis are treated alike once they no longer overlap. The Wasserstein distance, by contrast, charges for the displacement itself, so a near miss is scored as nearly correct. Figure 2.11 makes this contrast explicit by sliding and rescaling one spectrum against the other and tracking both measures.

A third family, *edit-distance* metrics, quantifies the minimum number of peak modifications needed to transform one spectrum into another, reflecting the structural differences between molecules. [49, pp. 1–2] Cosine similarity is also the conventional similarity metric used for library matching (Section 2.3.1), while the combined cosine–Wasserstein objective used to train the model of this work is developed in Section 3.3.6.

2.4 Computational Representations of Molecules and Reactions

The encoding of molecular structures and reactions in a machine-processable form is essential for the feasibility of making algorithmic predictions. [50, p. 518]

In computational chemistry, molecules can be represented as graphs. [50, p. 521] Reactions are modeled as graph transformations that break and form bonds. [51, p. 18]

2.4.1 Molecular Graphs

To enable algorithmic processing of molecular structures, we represent them using an abstract graph model that separates topology (which atoms are connected) from optional geometric detail. [51, p. 18]

Figure 2.12 illustrates the abstraction of a molecule as a graph, where vertices represent atoms and edges represent chemical bonds.




Figure 2.11: Spectrum A (red) is plotted upward and spectrum B (blue) mirrored downward on a shared m/z axis with non-uniformly spaced peaks.

a Aligned spectra: every B peak meets its A counterpart tip-to-tip, so cosine = 1 (intensity vectors parallel, $\theta = 0$) and Wasserstein = 0 (no mass to move).

b Shifting B by +8 Da slides every peak off its partner; cosine counts only intensity at *shared* positions, so its dot product vanishes and similarity collapses to 0, while Wasserstein measures the work to transport B's mass back onto A (amber arrows) and returns a finite 21.1. **c** Sweeping the shift, cosine (blue) drops to zero almost at once and stays flat—with *spurious* spikes near 14 and 17 Da where a displaced peak accidentally coincides with a *different* peak—whereas Wasserstein (amber) rises linearly with displacement, degrading gracefully.

d Scaling B's overall intensity with positions held fixed isolates the complementary property: cosine is unchanged at 1.00 because it compares vector *direction*, not magnitude (scale-invariant), while the (un-normalized) Wasserstein distance grows in proportion to the added mass. Together the panels show that the two measures capture different things: cosine rewards matching the *pattern* of relative intensities at fixed coordinates and ignores absolute scale, whereas Wasserstein is sensitive to peak *position* along the axis *and* to total intensity, making it robust to the small calibration shifts that break cosine but responsive to mass differences that cosine discards.




Figure 2.12: Toluene on the left depicted as a quantum-mechanical model with 3D coordinates and electron density, and on the right as a molecular graph abstraction. The graph captures the connectivity of atoms (vertices) and bonds (edges) while abstracting away geometric and electronic details.

This abstraction enables precise operations such as matching and rule-based transformation. [11, pp. 2–3] The formal definition used throughout this thesis is introduced next.

Formal definition

A molecular graph is typically defined as [52, p. 2]

$$G = (V, E) \tag{2.5}$$

where:

- V ... Vertices represent atoms
- E ... Edges represent chemical bonds between atoms

In Representations each node may carry attributes such as element type (e.g., C, N, O), formal charge, aromaticity, hybridization state, and isotope label. Likewise, each edge may encode bond order (single, double, or triple), aromatic character, and bond stereochemistry. [53, pp. 4–5] By construction, this basic molecular graph does not include geometric or electronic detail, such as 3D coordinates, molecular conformation, or wavefunction-level information; these require augmented graph models or separate geometric and electronic representations. [53, pp. 3–4] [54, p. 595] Embedding a molecule into 3D space, is itself a computationally hard problem, as it amounts to a high-dimensional conformational search over many degrees of freedom.

The abstraction is therefore a deliberate trade-off. What is lost are all properties that depend on spatial arrangement: stereochemistry and chirality, cis/trans and conformational isomerism, bond lengths and angles, and the electronic and energetic detail that determines reactivity,

spectra, and physical behaviour. Two molecules that differ only in handedness, for instance, map to the same underlying graph unless such information is explicitly re-encoded as vertex or edge attributes. [53, pp. 4–5] What is gained is a discrete, finite, and canonicalisable object that is amenable to exact algorithmic treatment: subgraph matching, isomorphism testing, and rule-based transformation all become well-defined combinatorial operations rather than continuous geometric or quantum-mechanical computations. [11, pp. 2–3] It is precisely this tractability that makes the graph representation the natural substrate for the transformation systems introduced next.

2.4.2 Chemical Graph Transformation Systems

To model chemical reactions algorithmically, molecular graphs are combined with rewrite rules that encode local bond changes. [55, p. 76] In this view, a reaction is not an arbitrary edit of a graph, but a constrained transformation that preserves valid molecular structure while modifying a specific reaction center. [11, p. 2] [56, p. 6]

Double Pushout Graph Rewriting

A standard formalism is the double pushout (DPO) approach. [11, p. 3] A rule is given by

$$p = (L \xleftarrow{l} K \xrightarrow{r} R),$$

where L is the left-hand side pattern to be matched, R is the right-hand side pattern after rewriting, and K is the interface that is preserved. [55, p. 75] Given a match $m : L \rightarrow G$ in a host graph G , DPO rewriting constructs an intermediate graph D (deletion pushout complement) [57, p. 680] and then the product graph H (insertion pushout). Intuitively, parts in $L \setminus K$ are removed, parts in $R \setminus K$ are added, and K remains unchanged. [11, p. 3] [58, p. 4]

$$\begin{array}{ccccc} L & \xleftarrow{l} & K & \xrightarrow{r} & R \\ m \downarrow & & \downarrow & & \downarrow m^* \\ G & \xleftarrow{l^*} & D & \xrightarrow{r^*} & H \end{array}$$

Figure 2.13: Schematic of the double pushout (DPO) graph rewriting process. A rule p consists of a left-hand side L , an interface K , and a right-hand side R . Given a match $m : L \rightarrow G$ in the host graph G , the DPO approach constructs an intermediate graph D by removing parts of G corresponding to $L \setminus K$, and then constructs the product graph H by adding parts corresponding to $R \setminus K$. The interface K is preserved throughout the transformation, ensuring that the resulting graph H maintains valid molecular structure.

In chemistry, this DPO construction (Figure 2.13) provides a representation of a reaction center: atoms and bonds in the preserved interface K define shared context, while changes between L and R encode bond breaking/forming events. [59, p. 5]

Derivation Graph

Building on the local rewrite step formalized above, derivation hypergraphs – induced by repeated rule application – provide a global view that shows mechanistically plausible chemical states. [11, pp. 2–3]

A derivation hypergraph, often also called a reaction network, is a higher-level hypergraph that captures all transformation steps generated by a graph rewriting system. [60, Sec. 10] [56, p. 2] Whereas a molecular graph represents the internal structure of a single molecule, a derivation hypergraph represents the *space of possible reactions* between many molecular states. [11, pp. 2–3]

Formally, each node corresponds to an object state (e.g., a molecule, multiset of molecules, or intermediate), and each directed edge corresponds to an application of a rewrite rule that transforms one state into another. [55, p. 75] If a rule p transforms state G into state H , this step is represented as

$$G \xrightarrow{p} H.$$

By collecting many such derivations, one obtains a directed hypergraph whose topology encodes reachable states, alternative pathways, and cycles.

In computational chemistry, derivation hypergraphs are useful for enumerating reaction pathways, identifying branching points, analyzing network connectivity, and ranking plausible mechanistic routes under additional constraints (e.g., energy, rule priorities, or experimental evidence). They therefore provide a natural bridge between rule-based chemistry models and downstream tasks such as pathway search, and machine-learning-based scoring of candidate transformations. [56, p. 2]

These graph-transformation and derivation hypergraph constructs are implemented in the MØD framework [59], [60], the software used throughout this work to generate fragmentation data; its Python module is referred to as `mod` in the following chapters.

One example derivation hypergraph for the fragmentation rules applied is in Figure 2.14.

Term Mode

Alongside the explicit graph representation, MØD can also handle molecular states in *term mode*, in which structures are expressed as formal terms (tree-like symbolic expressions) rather than as




Figure 2.14: Example derivation hypergraph showing reaction pathways for toluene fragmentation. Also illustrates the combinatorial explosion of possible pathways as rules are applied iteratively.

plain graphs. A term is built from function symbols (constructors such as `C`, `O`, or `bond`) applied to subterms, much like expressions in term-rewriting systems or functional programming, and is rewritten by the same rule machinery. The resulting object is a *term-graph*: a graph whose attributes carry this symbolic term structure. In this work term-graph inputs are normalised back to plain attributed graphs before encoding (Section 3.2.1), so term mode is relevant here only as the source format of some MØD outputs.

2.4.3 Graph Neural Networks for Molecular Representation Learning

The molecular encoder represents each molecule as a graph whose nodes correspond to atoms and whose edges correspond to chemical bonds. Initial atom features are updated through graph-convolutional message passing, allowing each atom representation to incorporate information from its local molecular environment. The resulting atom embeddings are pooled into a fixed-length graph-level representation, which is subsequently projected into the shared latent space used for comparison with spectral embeddings. This produces a continuous molecular vector that preserves structural information while being suitable for downstream neural-network training. For a visual overview of this workflow, see [Figure 2.15](#).




Figure 2.15: Starting from the molecular graph $G = (V, E)$ in **column 1**, where atoms form the nodes and bonds form the edges, each atom is initialized with a feature vector in **column 2** that encodes local chemical properties such as atom type, degree, formal charge, and aromaticity. Through repeated graph-convolution/message-passing layers in **column 3**, these atom representations are updated by aggregating information from neighboring atoms, yielding the context-aware atom embeddings shown in **column 4**. These embeddings are then pooled over all atoms in **column 5** to obtain a fixed-length molecular representation, which is projected in **column 6** into the shared latent space as the molecular embedding z_{mol} . As indicated in **column 7**, this vector can subsequently be compared with spectral embeddings for molecule-spectrum alignment. The atom features shown here (atom type, degree, formal charge, aromaticity) are a generic illustration of typical GNN inputs; the specific node-feature set used in this work—a one-hot element encoding together with formal charge, radical state, degree, aromaticity, and incident bond order—is defined in Section 3.2.1.

Chapter 3

Methodology

3.1 Datasets

This work draws from two main sources of input: measured mass spectra with their corresponding molecular structures as SMILES strings and a set of known fragmentation rules for the ionization method of interest.

3.1.1 Data Sources

Public Databases for spectra and structures

For this work, the National Institute of Standards and Technology (NIST) database for EI GC-MS was selected as the primary data source for the measured mass spectra. NIST is considered the gold standard in the field due to several key advantages: [61, p. 3]

- Comprehensive coverage of compounds
- Standardized data format ensuring consistency across entries
- Spectra experimentally measured under controlled, standardized instrument conditions

Alternative databases such as MassBank, PubChem, HMDB, and MoNa were considered but not utilized in this study.

The raw data from NIST consists of pairs of SMILES strings and their corresponding mass spectra.

Fragmentation Rules

The fragmentation reactions are defined by a set of rules derived from McLafferty’s seminal work on mass spectrometry fragmentation [10]. These were implemented as a rule extension to the




Figure 3.1: Schematic overview of the synthetic data-generation workflow. Starting from a sample molecule, fragmentation rules are applied to generate a fragmentation hypergraph that represents possible fragment formation pathways (left). The resulting fragments can also be recombined according to inverse rules to enumerate candidate molecules consistent with a given fragment set (right). Both directions are produced by the same rule-based graph-transformation pipeline, yielding the molecule–fragment data used to train the model.

MØD framework, enabling the generation of synthetic fragmentation data and the construction of derivation hypergraphs that capture the fragmentation pathways of molecules under study.

3.1.2 Preprocessing - Generating the Synthetic Data

Synthetic training data are generated with `mod` derivation hypergraphs (Figure 3.1). In this representation, nodes correspond to molecular states (fragments), while edges represent individual fragmentation transformations. The resulting hypergraph captures the reaction network and preserves fragment genealogy. This enables the tracing of parent-child relationships and the identification of reachable fragments. Subsequently, this information is utilised to validate predicted fragment assignments against chemically realistic pathways.

The `mod` strategy framework provides a systematic approach to explore chemical space by guiding fragmentation pathways through both physico-chemical properties and experimental data, such as mass spectra. A key advantage of this framework is its automatic generation of hypergraphs (reaction networks). This automation facilitates the construction of detailed atom flow networks for complex fragmentation patterns, thereby rendering it feasible to track chemical transformations and validate fragmentation predictions on a large scale. [59, p. 7]

Fragmentation Rule Application Strategies

The strategy layer accepts fragmentation rules as input and utilizes them to construct the corresponding reaction network by applying them to molecular structures. Listing 1 shows one representative rule definition used in this workflow.

```

1 IMS_4_9 = mod.Rule.fromDFS(
2     s =
3     "[C]1[C]2[C]3([C]4)[_A+.]5"
4     ">>"
5     "[C.]1" ". " "[C]2[C]3([C]4){=}[_A+.]5",
6     name =
7     "radical induced (alpha-)cleavage for a saturated site"
8     " 4.9"
9     " $R1R3R4Y5"
10 )

```

Listing 1: Sample Fragmentation Rule [10, p. 59]

In the `mod` framework, predicates are callback-based filters evaluated during rule application. More specifically, line 9 of Listing 1 is the critical component where the rule extension defines and enforces chemical constraints, removing candidate sites that violate them during the fragmentation process. This constraint specification forms the core mechanism of the rule extension, implemented as an extension to `mod` together with the broader application strategy.

Building on this mechanism, the rule extension serves as a *domain-specific language* (DSL) for precisely defining the spatial position and local chemical context of fragmentation sites. This positional language is not just a matter of convenience; it is essential for encoding complex chemical constraints and enables the systematic specification of reaction sites across molecular structures. Without it, encoding such constraints on a large scale would be impossible.

On the functional level, the decision flow of this filtering step, enabled by this rule extension, is summarized in Figure 3.2.

To operationalize these chemical constraints in the filtering stage, three predicate checks are




Figure 3.2: Flow diagram explaining predicate filter application at the fragmentation process. During rule application, candidate sites are generated based on the rule’s extension pattern. Each candidate is then evaluated by a series of predicate checks (heteroatom, saturation, alkyl) applied:

Heteroatom Check In order to ensure chemical plausibility during the application of the rule, the heteroatom check retains exclusively those candidate sites which contain at least one non-carbon, non-hydrogen atom within their local neighborhood.

McLafferty’s original rule definition [10, p. 57] does not specify which elements qualify as heteroatoms. Therefore, this work adopts the explicit design assumption that the admissible set comprises the organic elements B, C, N, O, P, S, F, Cl, Br and I, which is consistent with the typical scope of organic mass spectrometry. However, this assumption is not enforced programmatically, as the check accepts any non-carbon, non-hydrogen atom, regardless of the element. The resulting limitation is recorded in Section 3.2.5.

The heteroatom check evaluates the neighborhood of the rule-extension location via local graph traversal. A candidate is accepted only if at least one non-carbon, non-hydrogen atom is found, formulated as $\text{BFS} \setminus \{\text{H}, \text{C}\} \neq \emptyset$. If the explored neighborhood contains only carbon and hydrogen, the candidate is rejected. Figure 3.3 contrasts both cases: a valid neighborhood containing a heteroatom and an invalid neighborhood without one.




Figure 3.3: Heteroatom check predicate: valid case with detected heteroatom (left) versus invalid C/H-only neighborhood (right).

Saturation Check The saturation check is formulated as a constrained path search between a start and an end atom. A candidate path is accepted only if it follows a carbon backbone through single C–C bonds, while side branches are limited to hydrogen. This removes candidates that include incompatible branch chemistry, such as additional carbon side branches. Figure 3.4 highlights the accepted saturated path in green and marks a rejected carbon branch.




Figure 3.4: Illustration of the concept of a saturated path in a molecular graph. A valid saturated path connects the start and end atoms through a carbon backbone consisting exclusively of single C–C bonds, with only hydrogen atoms as side branches (top). In contrast, a path is not considered saturated if additional carbon side branches are attached to the backbone (bottom).

Alkyl Check The alkyl check ensures that the fragmentation site is part of an alkyl group, restricting fragmentation to aliphatic carbon chains, where such reactions are most likely to occur. This constraint helps filter out fragmentation events in aromatic rings or other structural contexts where different chemical reactivity dominates.

Figure 3.5 visualizes the alkyl check as a breadth-first traversal of the local neighborhood around the rule-extension location. Unlike the saturation check, which validates a *path* between a designated start and end atom, the alkyl check inspects the *surrounding* atoms and accepts the site only if every reachable neighbor is a carbon or hydrogen atom. The presence of any heteroatom, such as the nitrogen marked on the right of Figure 3.5, causes the predicate to fail, since the site no longer lies within a purely aliphatic (alkyl) environment.




Figure 3.5: Alkyl check predicate formulated as a local-neighborhood traversal. Starting from the rule extension location, a breadth-first search (BFS) is performed over the local molecular graph to verify that all reachable neighboring atoms belong to the set $\{C, H\}$. A matched substructure is accepted as an alkyl group if only carbon and hydrogen atoms are encountered (left), whereas the presence of a heteroatom such as nitrogen causes the predicate to fail (right).

3.1.3 Data Splits

The dataset was partitioned in two stages into training, validation, and test subsets. In the *first* stage, the full set of molecule-spectrum pairs was divided into an 80% training portion and a 20% test portion. In the *second* stage, the initial training portion was further split into 70% training and 30% validation. This results in a final distribution of 56% training data, 24% validation data, and 20% test data relative to the original dataset.

These ratios were chosen as a pragmatic trade-off between model development and robust evaluation. The 80/20 split keeps a sufficiently large hold-out test set for a meaningful final performance estimate, while still reserving most samples for development.

This splitting strategy serves two purposes. *First*, it preserves a dedicated hold-out test set

that is never used during model development and is therefore reserved for the final evaluation. *Second*, it provides a separate validation set for model selection, hyperparameter tuning, and intermediate performance monitoring. Having distinct, non-overlapping training, validation, and test sets ensures that no information from the data used for evaluation influences model fitting, thereby preventing information leakage.

The splits are made with `shuffle=False`, so the data are partitioned in their original order rather than being randomly permuted beforehand. The split is therefore deterministic and reproducible by construction; no random permutation is involved, and the `random_state` setting has no effect in this configuration.

3.2 Latent Representation Schemes

In this work, we encode molecular fragmentation data through three tightly coupled representation streams:

1. **Molecular states** as chemistry-aware graphs
2. **Derivation dynamics** as a hypergraph over molecular states and transformation rules
3. **Mass spectra** as peak-intensity vectors aligned with structure-derived latent states

The core design principle is a two-level encoding pipeline. *First*, each molecular state is encoded at the level of atoms and bonds. *Second*, these state-level encodings are used as node features in a derivation hypergraph. This approach preserves local chemical structure while explicitly representing multi-step fragmentation logic. Figure 3.6 provides a schematic overview of this pipeline.



Figure 3.6: Overview of the two-level molecular encoding pipeline: local graph encoders capture atom- and bond-level chemistry, while global hypergraph encoders model multi-step fragmentation over embedded molecular states.

In particular, a molecule is not encoded as a SMILES sequence, fingerprint, or direct spectrum-oriented tokenization. Instead, it is represented as an explicit graph object derived from the MØD framework, then mapped into a dense latent embedding, and finally placed into a derivation-hypergraph context together with learned rule representations.

3.2.1 Molecular Graph Encoding

PyTorch Geometric (PyG) [62] is a library that extends PyTorch to enable efficient processing of graph-structured data. It provides specialized data structures and operations optimized for graph neural networks (GNNs).

For molecular graph encoding, each molecular state is converted into a PyG graph object following the molecular-graph definition of Section 2.4.1, so that GNN architectures can learn from both node features (atomic properties) and edge features (bond information).

Node features encode both atom identity and the local connectivity that enables the encoder to distinguish between constitutional isomers rather than collapsing to composition (and therefore mass). For each atom, the feature vector concatenates a one-hot element encoding over a vocabulary with five scalar channels:

- formal charge (**charge**), the atom’s formal charge
- radical state (**radical**), whether the atom carries an unpaired electron
- degree (**degree**), the number of incident bonds
- aromaticity (**aromatic**), a flag set on atoms that lie in an aromatic bond
- incident bond order (**bond_order**), the summed order of the atom’s bonds

Edge features encode bond type as one scalar value, with single, double, triple, and aromatic bonds mapped to 1.0, 2.0, 3.0, and 1.5, respectively. Connectivity is stored in sparse Coordinate (COO) format, where edges are specified as pairs of row and column indices, enabling efficient storage and computation for graphs with varying density.

Before featurization, term-graph inputs are detected and normalized to standard graph form, so the encoder always receives a consistently formatted representation.

From these layers we use the edge-aware **GINEConv**, a graph-isomorphism-network variant that consumes the bond-order edge features alongside the atom features, to embed each molecular graph into the latent space (detailed in Section 3.3.3).

3.2.2 Two-Level Molecular Latent Representation

After graph featurization, each molecular graph is passed to the previously Section 3.2.1 described encoder, which compresses it into a dense state embedding, one vector per molecular state.

The central, unique architectural feature is that molecules are not treated as isolated entities. Each molecular state corresponds to a vertex in the *derivation hypergraph*, the functional object

that records fragment genealogy. For computation, the derivation hypergraph is realised as a hypergraph data structure, in which a single hyperedge can connect all fragment vertices produced by one transformation step. Throughout this work, "derivation hypergraph" refers to this functional object, whereas "hypergraph" denotes only its concrete data-structure encoding. Encoding each state graph in turn yields one embedding per vertex, so that the node features of the derivation hypergraph are not raw atom features but learned embeddings of complete molecular states.

This yields two explicit representational levels:

1. Lower level: each molecule is a molecular graph (atoms and bonds).
2. Higher level: each molecular state becomes a vertex in the derivation hypergraph.

Molecule encoding therefore proceeds in three steps:

1. Chemistry-aware graph feature extraction
2. Graph-to-vector compression into a dense state embedding
3. Use of this embedding as the vertex feature in the derivation hypergraph

In the resulting derivation hypergraph data structure, each vertex carries its state embedding as a node feature. The incidence structure records which vertices participate in each hyperedge. Every hyperedge is annotated with a learned embedding of the transformation rule (or set of rules) that produced the corresponding derivation step (Section 3.2.3). This two-level construction is illustrated in Figure 3.7.




Figure 3.7: In the *first* stage, each molecular state is represented as an atom-bond graph and encoded by a learned graph encoder into a dense state embedding. In the *second* stage, these state embeddings are integrated into the derivation hypergraph as vertex features. Thus, vertices correspond to molecular states, while hyperedges and incidence relations encode the higher-level fragmentation structure.

3.2.3 Rule Representation and Hyperedge Annotation

Rule processing is handled separately from molecule processing. Unlike a molecular state, a rule is not expanded into a chemically rich multi-node graph; instead, each rule is represented compactly as a single-node graph whose only feature is its indexed rule identifier.

Rule handling then proceeds in three stages: collection, encoding and aggregation. If there are several rules associated with one hyperedge, their embeddings are mean-pooled to create a single hyperedge-level vector. Consequently, each hyperedge is annotated with a learned representation of the transformation rule (or set of rules) that produced that derivation step.

3.2.4 Spectral Vector and Embedding Representation

Each mass spectrum is treated as a set of tuples $(m/z, I)$, where m/z is the mass-to-charge ratio and I is the corresponding peak intensity. For model input, these tuples are converted into peak-intensity vectors over m/z bins. This tuple-based encoding preserves peak information while producing a dense, differentiable input for neural encoders. These vectors are embedded

into a spectral latent space, which is trained to be compatible with the structure-derived latent representation. A shared latent view integrates fragmentation topology, molecular-state embeddings, and observed spectral evidence in downstream learning tasks.

The conversion of a peak list into a latent embedding z_{spec} proceeds in three steps (Algorithm 1). *First*, the variable-length tuple list is *binned* into a fixed-length vector by discretizing the m/z axis into B equal bins and accumulating peak intensities per bin, yielding an input whose dimension is independent of the number of peaks. *Second*, an elementwise square root *compresses* the wide intensity range, and L_2 normalization removes overall-scale differences so that the encoder responds to the relative-intensity pattern rather than absolute magnitude. *Third*, the trained encoder Enc_{spec} maps this vector to $z_{\text{spec}} \in \mathbb{R}^d$, the spectral embedding that the contrastive objective (Equation (3.4)) aligns with the structure-derived representation.

3.2.5 Disadvantages

Despite representational advantages, the proposed encoding pipeline also introduces several methodological and practical limitations:

- Loss of information due to dimensionality reduction from molecular graphs to fixed-size latent vectors
- Limited interpretability of latent dimensions and pooled rule embeddings
- Potential propagation of preprocessing artifacts into learned representations
- Risk of over-squashing and long-range dependency loss in message-passing encoders
- The set of admissible heteroatoms is not enforced in code but treated as an implicit design assumption (see Section 3.1.2), so structures containing elements outside the assumed organic set are not formally excluded
- The MØD backend used for data generation does not support biradicals, so fragmentation pathways through species with two unpaired electrons cannot be represented

3.3 Model Architecture

This section states the learning problem, then describes the shared three-modality architecture and its forward (structure-to-spectrum) and backward (spectrum-to-structure) branches.

Algorithm 1 Spectrum to latent embedding: from $(m/z, \text{intensity})$ tuples to z_{spec}

Require: peaks: list of (mz, I) tuples; mz_min , mz_max , bin_width ; trained Enc_{spec}

Ensure: latent embedding $z_{\text{spec}} \in \mathbb{R}^d$

```

1: // Step 1: bin the sparse peak list into a dense vector
2:  $B \leftarrow \lfloor (mz\_max - mz\_min) / bin\_width \rfloor$  ▷ number of bins
3:  $vec \leftarrow \mathbf{0}_B$ 
4: for all  $(mz, I)$  in peaks do
5:    $idx \leftarrow \lfloor (mz - mz\_min) / bin\_width \rfloor$ 
6:   if  $0 \leq idx < B$  then
7:      $vec[idx] \leftarrow vec[idx] + I$  ▷ accumulate intensity
8:   end if
9: end for

10: // Step 2: dynamic-range compression + normalization
11:  $vec \leftarrow \sqrt{\max(vec, 0)}$  ▷ elementwise
12: if  $\|vec\|_2 > 0$  then
13:    $vec \leftarrow vec / \|vec\|_2$  ▷ L2 normalize
14: end if

15: // Step 3: encode into latent space; equivalent to nn.Sequential
16:  $h \leftarrow \text{MLP}_1(vec)$  ▷ "hidden" representation
17:  $z_{\text{spec}} \leftarrow \text{MLP}_2(h)$ 
18: return  $z_{\text{spec}}$ 

```

3.3.1 Problem Formulation

Each training example couples three views of the same molecule:

- a *molecular graph* $G = (V, E)$ (Section 2.4.1) representing the precursor structure;
- a *derivation graph* $D(G)$ whose vertices are fragment states and whose edges encode the fragmentation steps that connect them;
- a measured *mass spectrum*, binned into a fixed-length intensity vector $s \in \mathbb{R}^B$ over B equal m/z bins (Section 3.2.4).

Three encoders map these views into latent vectors,

$$z_{\text{mol}} = \text{Enc}_{\text{mol}}(G), \quad z_{\text{frag}} = \text{Enc}_{\text{frag}}(D(G)), \quad z_{\text{spec}} = \text{Enc}_{\text{spec}}(s), \quad (3.1)$$

where all three encoders produce embeddings of the same dimension d , and a set of task heads projects these views into the shared latent space $\mathbb{R}^d$ in which structural and spectral embeddings can be compared directly.

The model is then trained in two coupled directions:

Forward (structure $\rightarrow$ spectrum): the molecular and fragment embeddings are fused into a combined latent z_{combined} and decoded into a predicted spectrum $\hat{s} = \text{Dec}_{\text{spec}}(z_{\text{combined}})$.

Backward (spectrum $\rightarrow$ structure): the spectral embedding z_{spec} is aligned with the molecular embedding z_{mol} , so that at inference time the molecules whose embeddings lie closest to a query spectrum can be retrieved.

Both directions are optimised jointly under a single composite objective $\mathcal{L}_{\text{total}}$ (Section 3.3.6); structure-to-spectrum prediction and spectrum-to-structure retrieval therefore share, and mutually constrain, the same latent space. Table 3.1 collects the symbols used throughout this section.

Table 3.1: Notation used in the model description.

Symbol	Meaning
G	molecular graph (atoms as nodes, bonds as edges)
$D(G)$	derivation graph of fragment states for molecule G
$s \in \mathbb{R}^B$	binned mass spectrum over B m/z bins
$\mathbf{h}_{s_i} \in \mathbb{R}^d$	embedding of an individual molecular state s_i
z_{mol}	molecular (precursor) latent
z_{frag}	fragment / derivation-graph latent
z_{combined}	fused forward latent, $\frac{1}{2} (z_{\text{mol}} + z_{\text{frag}})$
z_{spec}	spectral latent
z_{fwd}	forward-head projection of z_{combined} (decoder conditioning, Phase A)
z_{bwd}	backward-head projection of z_{spec} (alignment / retrieval query, Phase B)
$\hat{s}$	predicted spectrum
$\text{Enc}_{\text{mol}}, \text{Enc}_{\text{frag}}, \text{Enc}_{\text{spec}}$	molecular / fragment / spectral encoders
Dec_{spec}	fragment-grounded spectrum decoder
$\text{Head}_{\text{FWD}}, \text{Head}_{\text{BWD}}$	forward / backward task-head projections
d	shared latent dimension (graph embeddings, alignment / retrieval)

3.3.2 Overview of the Learning Framework

The methodology presented in this chapter constitutes the central original contribution of this work: to the author’s knowledge, it is the first framework to couple MØD-derived fragmentation graphs with a learned two-level molecular representation and to align spectral and structural embeddings in a single bidirectional model. Both the rule-extension DSL and the coupled forward–backward training objective are developed here for the first time.

The implemented learning framework is a bidirectional model that couples molecular structure, fragment-level information, and mass spectra in a shared latent space. It operates on three synchronized views of the same molecule, each handled by a dedicated encoder that ap-

pears along the top of Figure 3.8: a molecular graph, embedded by a graph neural network, realising Enc_{mol} ; a derivation tree of fragment states with their masses, encoded by a backend that represents the fragment graphs and their parent-child relations explicitly, realising Enc_{frag} ; and a binned mass spectrum, projected into the latent space by a dedicated spectrum encoder, realising Enc_{spec} ; see Section 3.2.4. All three encoders emit vectors of the same dimension d , so their embeddings share one comparable space that supports both forward spectrum prediction and backward structure retrieval.

A central design decision is that the model is *not* a direct spectrum-to-SMILES generator but a bidirectional embedding model. In the forward direction, spectra are predicted from molecular and fragment information. In the backward direction, a spectrum is mapped into the shared latent space and the molecular embeddings that lie closest to it are retrieved. Therefore, the inverse path should not be considered a decoder that explicitly reconstructs a molecular graph, but rather a nearest-neighbour retrieval step over a precomputed index of known structures, as shown in Figure 3.9.

As it is not possible to show both how these components are wired together for training and how the backward branch runs independently at inference time in a single diagram, the architecture is presented through two complementary figures that share the same encoder blocks and latent symbols. Figure 3.8 gives the training dataflow. All three encoders, the latent fusion, the shared task heads, and the fragment-grounded decoder, with the loss edges that couple the two directions. Figure 3.9 then isolates the inference pathway, in which a query spectrum is embedded and used to retrieve and rerank candidate molecules. The subsections that follow trace each branch against these two figures.




Figure 3.8: Training-time dataflow of the bidirectional framework. *Phase A* (blue): a molecule graph and its derivation tree are encoded to z_{mol} and z_{frag} , averaged into $z_{\text{combined}} = (z_{\text{mol}} + z_{\text{frag}})/2$, projected by Head_{FWD} , and decoded to $\hat{s}$ by the shared fragment-grounded decoder, which places one intensity per fragment at that fragment’s mass. *Phase B* (amber): the binned spectrum is encoded to z_{spec} and projected by Head_{BWD} to z_{bwd} , which is aligned to the molecular latent z_{mol} by an InfoNCE objective and conditions the same decoder. Dotted edges denote spectrum-reconstruction losses (cosine and Wasserstein); the full set of loss terms is listed in Table 3.2.




Figure 3.9: Inference-time spectrum-to-structure retrieval. A query spectrum is encoded to z_{bwd} through Enc_{Spec} and Head_{BWD} and used as a nearest-neighbour query against the molecular index—a precomputed store of normalized molecular latents z_{mol} (*build*). The retrieved *top-k* candidates are then reranked by the forward model, which predicts each candidate’s spectrum and scores it against the observed input, yielding the final ranked structures.

3.3.3 Forward Model: Structure-to-Spectrum Prediction

The forward branch is the Phase A pathway in Figure 3.8, running from the **Molecule** and **Derivation tree** inputs down to the shared decoder. Its purpose is to predict a mass spectrum from a molecular structure together with fragment-level information derived from precomputed derivation hypergraphs. The molecular encoder Enc_{mol} is implemented as an edge-aware graph neural network built from GINEConv layers (a graph-isomorphism-network variant that incorporates edge features), which consume the bond-order edge attributes together with the atom features. A concatenated mean-and-max readout then maps the molecular graph to a latent vector z_{mol} ; retaining the max component helps preserve the few local atom environments that distinguish constitutional isomers, which a mean-only pooling would dilute. This corresponds to the molecular branch shown in the figure.

The fragment encoder is the **FragSetEncoder** acting on the **Derivation tree** input in Fig-

ure 3.8, and encodes derivation hypergraph structure explicitly. For each molecule, the dataset constructs fragment graphs and an explicit derivation hypergraph whose nodes correspond to fragments and whose directed edges encode fragmentation relations. The fragment encoder first embeds the individual fragment graphs using the molecular graph encoder as a backbone. The fragment masses are added as extra features for the nodes, and then a second graph neural network is applied to the derivation hypergraph itself. A global pooling step yields a single latent fragment representation z_{frag} , while the per-fragment node embeddings and masses are retained separately, the `node_embs` and `node_masses` tensors in Figure 3.8, for later use by the fragment-grounded decoder. Accordingly, this branch is implemented as a graph encoder over derivation hypergraphs rather than as a simple pooling operation over an unordered set of fragments.

The latent-fusion node drawn as $z_{\text{combined}} = (z_{\text{mol}} + z_{\text{frag}})/2$ in Figure 3.8 is implemented by a simple arithmetic average of the molecular and fragment latents,

$$z_{\text{combined}} = \frac{z_{\text{mol}} + z_{\text{frag}}}{2}. \quad (3.2)$$

This fused representation is then passed through the forward task head, producing z_{fwd} , which projects the latent vector into a task-specific subspace.

In latent fusion, the molecular and fragment branches are designed to produce vectors of the same dimensionality, either by sharing the final latent dimension d or by applying linear projection layers in their respective task heads. This ensures z_{mol} and z_{frag} lie in the same vector space and can be combined element-wise. The simple arithmetic average used here Equation (3.2) is a lightweight fusion strategy that treats both modalities equally; more complex fusion (concatenation followed by a projection, attention-based fusion, etc.) could be used.

Finally, the spectrum decoder turns this representation into a predicted spectrum $\hat{s}$. Rather than mapping the fused latent to a free-form intensity vector over all m/z bins, the decoder is *fragment-grounded*: for each fragment it predicts a single intensity from that fragment’s node embedding and mass (the `node_embs` and `node_masses` inputs to the decoder in the figure) together with the forward conditioning latent z_{fwd} , and then scatters these intensities onto the m/z bins fixed by the fragments’ own masses. Therefore, the predicted spectrum is a function of the actual fragment set. By construction, every peak lands at a physically realizable fragment mass, and the decoder cannot collapse to a single, input-independent spectrum. This design is deliberate: a free-form decoder that regresses the full binned intensity vector directly from the pooled latent would collapse onto an input-independent spectrum, a failure mode that the fragment-grounded head removes by construction and that the results chapter examines (Chapter 4).

3.3.4 Backward Model: Spectrum-to-Structure Retrieval

The backward branch is the Phase B (amber) pathway in Figure 3.8: the **Binned spectrum** input, the encoder, and the backward head. Its function is to encode spectra into a latent representation that can be aligned with molecular embeddings and used for later retrieval. The spectrum encoder receives a fixed-length vector of binned intensities. It is implemented as a regularised stack of residual multi-layer perceptron blocks incorporating normalisation and dropout. This makes the spectral branch deeper and more expressive than the single **EncSpec** box in the figure suggests.

Given an input spectrum, produces a latent vector z_{spec} , which is then passed through the backward task head to give the projected embedding z_{bwd} . This head serves two purposes. *First*, it provides a task-specific projection for latent alignment with molecular embeddings during training. *Second*, it defines the representation that is used at inference time for inverse retrieval. The projected spectrum embedding is then compared to a latent index of molecular embeddings, thus enabling the model to retrieve the most similar candidate structures in embedding space.

At inference time only this branch runs, and it becomes the spectrum-to-latent-to-candidate pipeline of Figure 3.9: z_{bwd} is issued as a nearest-neighbour *query* against a **LatentIndex** built from a precomputed, normalized database of molecular latents z_{mol} , returning the top- k candidate molecules by cosine similarity; these are then optionally reranked by passing each candidate through the forward branch and scoring its reconstructed spectrum against the query (the **Forward-model rerank** box).

3.3.5 Bidirectionality

The two directions are not learned independently. As Figure 3.8 shows, both feed the *same* fragment-grounded decoder, the **DecSpecFragment** box receives the forward conditioning z_{fwd} (Phase A) and the backward conditioning z_{bwd} (Phase B) and, because both branches are trained jointly, information learned in one direction constrains the other. At inference time the molecular encoder builds an index over known candidate structures while the spectrum encoder provides the query, so forward prediction and inverse retrieval become two views of the same embedding problem. The forward model further feeds back into retrieval through the reranking step of Figure 3.9, where retrieved candidates are rescored by how well their forward-predicted spectra match the observed spectrum.

3.3.6 Loss Function

The forward model is trained against a differentiable spectral measure that also serves as the objective, meaning its geometry shapes the loss landscape. The geometric *cosine similarity* and the transport-based *Wasserstein distance* are complementary (Section 2.3.5), with cosine rewarding matching of relative intensities but ignoring peak displacement, and Wasserstein charging for displacement and degrading gracefully. We therefore combine them into the reconstruction loss.

$$\mathcal{L}_{\text{spec}} = \alpha_{\text{cos}} (1 - \cos(P, Q)) + \lambda \text{WD}(P, Q), \quad (3.3)$$

where $1 - \cos(P, Q)$ expresses the cosine as a dissimilarity minimised at agreement, and the weights $\alpha_{\text{cos}} = 1.59$ and $\lambda = 1.17$ (Table A.1) balance intensity-pattern fidelity against positional accuracy, with the cosine term evaluated on the precursor-mass-masked prediction.

This reconstruction term is only the anchor of a multi-component objective that adds chemically motivated constraints and a contrastive alignment term; the resulting *loss shaping* stabilises training and ties the shared latent space to chemical structure, with Table 3.2 listing every term by branch and weight and eq. (3.8) assembling the full objective. On the forward branch, $\mathcal{L}_{\text{spec}}$ is the regularisation of the spectrum is achieved through a penalty for regions that are forbidden, which suppresses intensity beyond the precursor mass. There is also a cycle term that re-encodes the predicted spectrum towards the molecular latent state, as well as a diversity term that discourages near-identical predictions within a batch. The latter is now a soft regulariser rather than a defence against collapse, since the fragment-grounded decoder (see Section 3.3.3) mmakes each predicted spectrum an explicit function of its molecule’s fragment set. This means that distinct inputs already yield distinct outputs by construction. The backward branch reuses the reconstruction and forbidden-region terms, adding a spectrum-reconstruction cycle and contrastive alignment between the spectral and molecular embeddings. This contrastive term serves as both the reconstruction objective and the retrieval-ranking objective because batches are grouped by precursor mass. This means that the negatives in each batch are nearly isobaric. The retrieval task is to rank the true molecule above them.

Two building blocks recur across these terms. The first is the InfoNCE alignment objective, a contrastive loss that pulls matching pairs together in the latent space while pushing non-matching pairs apart; for a batch of N matched pairs $(z_{\text{spec},i}, z_{\text{mol},i})_{i=1}^N$,

$$\mathcal{L}_{\text{InfoNCE}} = -\frac{1}{N} \sum_{i=1}^N \log \frac{e^{\frac{\text{sim}(z_{\text{spec},i}, z_{\text{mol},i})}{\tau}}}{\sum_{j=1}^N e^{\frac{\text{sim}(z_{\text{spec},i}, z_{\text{mol},j})}{\tau}}}, \quad (3.4)$$

where $\text{sim}(\cdot, \cdot)$ is cosine similarity and τ the temperature; the denominator turns alignment into an N -way classification in which the correct molecule $z_{\text{mol},i}$ must be ranked above the

other in-batch embeddings, which serve as negatives without separate mining, and the two contrastive terms in Table 3.2 instantiate this form with different embedding pairs. The second is a homoscedastic task-uncertainty weighting that wraps selected terms,

$$\mathcal{U}(\mathcal{L}; \rho) = e^{-\rho} \mathcal{L} + \rho, \quad (3.5)$$

where $\rho = \log \sigma^2$ is learned (initialised to 0): $e^{-\rho} = 1/\sigma^2$ down-weights a noisy term while the additive ρ prevents the degenerate solution $\sigma \rightarrow \infty$.

Table 3.2: Loss terms with the branch in which each is applied and its weight. The reconstruction components $\mathcal{L}_{\text{cos}}$ and $\mathcal{L}_{\text{wass}}$ together form $\mathcal{L}_{\text{spec}}$ (eq. (3.3)), which is wrapped in the uncertainty term $\mathcal{U}(\cdot; \rho_{\text{spec}})$. Here z_{bwd} is the backward-head projection of z_{spec} and m the hard precursor-mass mask.

Term	Compares	Branch	Weight
$\mathcal{L}_{\text{cos}}$	masked $\hat{s} \odot m$ vs. s (cosine)	both	$\alpha_{\text{cos}} = 1.59$
$\mathcal{L}_{\text{wass}}$	$\hat{s}$ vs. s (1-D EMD over bins)	both	$\lambda = 1.17$
$\mathcal{L}_{\text{forb}}$	intensity beyond precursor mask	both	$\beta = 0.1$
$\mathcal{L}_{\text{cyc}}$	$\text{Enc}_{\text{spec}}(\hat{s})$ vs. z_{mol}	forward	$\lambda_{\text{cyc}} = 0.77$
$\mathcal{L}_{\text{div}}$	off-diagonal batch correlations	forward	$\alpha_{\text{div}} = 0.1$
$\mathcal{L}_{\text{cyc-spec}}$	$\text{Enc}_{\text{spec}}(\hat{s})$ vs. z_{spec} (cosine cycle)	backward	$\lambda_{\text{cyc-spec}} = 0.089$
$\mathcal{L}_{\text{con}}$	z_{bwd} vs. z_{mol} (InfoNCE)	backward	$\lambda_{\text{con}} = 0.1$
$\mathcal{L}_{\text{align-sm}}$	z_{spec} vs. z_{mol} (InfoNCE alignment)	backward	$\lambda_{\text{align}} = 0.30$

The model is trained in two phases, each with its own optimiser step. The forward phase minimises

$$\mathcal{L}_{\text{fwd}} = \mathcal{U}(\mathcal{L}_{\text{spec}}; \rho_{\text{spec}}) + \beta \mathcal{L}_{\text{forb}} + \lambda_{\text{cyc}} \mathcal{L}_{\text{cyc}} + \alpha_{\text{div}} \mathcal{L}_{\text{div}}, \quad (3.6)$$

and the backward phase minimises

$$\begin{aligned} \mathcal{L}_{\text{bwd}} = & \mathcal{U}(\mathcal{L}_{\text{spec}}; \rho_{\text{spec}}) + \lambda_{\text{con}} \mathcal{U}(\mathcal{L}_{\text{con}}; \rho_{\text{con}}) + \beta \mathcal{L}_{\text{forb}} \\ & + \lambda_{\text{cyc-spec}} \mathcal{L}_{\text{cyc-spec}} + \lambda_{\text{align}} \mathcal{L}_{\text{align-sm}}. \end{aligned} \quad (3.7)$$

Over a full epoch the model therefore optimises the union of these terms,

$$\begin{aligned} \mathcal{L}_{\text{total}} = & \mathcal{U}(\mathcal{L}_{\text{spec}}; \rho_{\text{spec}}) + \beta \mathcal{L}_{\text{forb}} + \lambda_{\text{cyc}} \mathcal{L}_{\text{cyc}} + \lambda_{\text{cyc-spec}} \mathcal{L}_{\text{cyc-spec}} + \lambda_{\text{align}} \mathcal{L}_{\text{align-sm}} \\ & + \lambda_{\text{con}} \mathcal{U}(\mathcal{L}_{\text{con}}; \rho_{\text{con}}) + \alpha_{\text{div}} \mathcal{L}_{\text{div}}, \end{aligned} \quad (3.8)$$

with $\alpha_{\text{cos}} = 1.59$, $\lambda = 1.17$, $\lambda_{\text{cyc}} = 0.77$, $\lambda_{\text{cyc-spec}} = 0.089$, and $\lambda_{\text{align}} = 0.30$ selected by the hyperparameter search (Table A.1), while $\beta = 0.1$, $\lambda_{\text{con}} = 0.1$, $\alpha_{\text{div}} = 0.1$, and $\tau = 0.07$ were fixed a priori; the log-variances ρ_{spec} and ρ_{con} are learned.

There are two things to bear in mind. *First*, the reconstruction, forbidden-region and cycle terms enter *both* phases, so eq. (3.8) is the union of the two phase objectives rather than a single backward pass. *Second*, the weighting U adapts the balance between reconstruction and contrastive terms during training rather than fixing it by hand. Therefore, the objective is not simply encoder-decoder reconstruction, but rather a coupled bidirectional system that jointly optimises reconstruction, alignment, ranking and chemical-plausibility constraints over the shared latent space.

3.3.7 Training Procedure

The two branches (forward and backward) are trained *sequentially*, in two separate phases, rather than jointly or alternated per batch. Phase A trains the forward branch (structure $\rightarrow$ spectrum) to convergence. After this, Phase B trains the backward branch (spectrum $\rightarrow$ structure alignment and spectrum reconstruction). Each phase runs for a fixed epoch budget under its own Adam optimiser.

Within an epoch, the loss terms of the active branch are summed and back-propagated in a single optimiser step per batch. The spectrum decoder and the task heads are shared across both phases and are therefore updated by both training optimisers. As the molecule encoder is a submodule of the fragment encoder, it is also updated in Phase A, even though Phase A’s optimiser is constructed over the fragment encoder, the spectrum decoder and the task heads.

Both phases use identical Adam settings, with no learning-rate schedule, warmup, or gradient clipping. Training is purely epoch-budgeted and the final-epoch weights are retained as the trained model. Accordingly, the test split is used only *post hoc* for retrieval evaluation (Chapter 4), never for checkpoint selection. Reproducibility rests on a single global seed for initialisation and a fixed split state (Section 3.1.3). The full set of hyperparameters is collected in Table A.1.

Chapter 4

Results and Discussion

4.1 Computation Times of Fragmentation Expansion

Computation times present a significant bottleneck for this approach. The `mod` framework, while powerful in generating comprehensive reaction networks, prioritizes flexibility and completeness over computational efficiency. Combined with the rapid growth in the number of possible fragments and parent molecules, runtime becomes prohibitive as the number of fragmentation rules increases. Since only a subset of known rules were applied in this study, computation times per molecule will rise substantially as additional rules are incorporated, making further optimization necessary to ensure the feasibility of this approach at scale.

4.1.1 Visualization of Computation Times

A regression analysis was performed to model the relationship between molecular complexity and computation time.

Here, molecular complexity refers to Boettcher’s complexity measure [63], which summarizes structural intricacy from the molecular graph. This provides a single scalar descriptor that enables molecules of different sizes and connectivities to be compared on a common scale. In this analysis, it is used as an indicator of fragmentation difficulty. However, runtime is also influenced by factors such as the molecule’s specific structure and the available fragmentation pathways.

The analysis assumes:

- Homoscedasticity of residuals
- Independence of residual observations
- Normality of residual distribution

In Figure 4.1, the regression results are visualized, showing a positive correlation between molecular complexity and computation time, but also show a large variance in computation times for molecules of similar complexity. Note that the threshold of 40 seconds was chosen arbitrarily to exclude trivial expansions.




Figure 4.1: Regression of Computation time, excluding molecules that had trivial expansions (time < 40s). The plot shows a positive correlation between molecular complexity and computation time, but also a large variance in computation times for molecules of similar complexity.

The large variance suggests that other factors, such as specific structural features or the particular fragmentation pathways available, also play a significant role in determining computation time.

4.2 Evaluation of Structure–Spectrum Translation

The evaluation must control for one confound above all others: the precursor (parent) mass. Because a candidate whose mass matches the query’s implied precursor mass is already a strong match, any retrieval score that is not held against a mass baseline can look successful while doing little more than mass filtering.

A free-form spectrum decoder can collapse onto the dataset-mean spectrum, predicting essentially the same intensity vector regardless of the input molecule, while a reranking stage built on top of it still appears to improve retrieval substantially. Such an apparent gain is almost

entirely due to a precursor-mass consistency filter: replacing the learned forward spectrum with a constant mean spectrum, or even an all-ones vector, while keeping the mass mask reproduces the same reranking benefit, so the learned spectral shape itself contributes little.

First, the fragment-grounded decoder, Section 3.3.3 makes each predicted spectrum an explicit function of the molecule’s fragment set. This means that it cannot collapse to an input-independent spectrum, and that its peaks will land at physically realizable fragment masses by construction. *Secondly*, retrieval is reported under a *mass-controlled* scheme: within a fixed window around the query’s true mass, the learned encoder retrieval $\cos(z_{\text{spec}}, z_{\text{mol}})$ is compared against the analytic random-chance baseline for a pool of that size. Only a result that beats chance *inside* the window counts as learned structural signal, because within a narrow mass window the mass is approximately constant and can no longer explain the ranking.

4.2.1 The Forward Model Is Molecule-Specific

With the fragment-based decoder, forward prediction is input-dependent by construction. Since each predicted spectrum is explicitly dependent on the molecule’s own set of fragments, the predictions vary according to the input rather than reproducing a single average spectrum. For a representative molecule such as benzene, the predicted peaks fall at fragment masses $m/z \approx 26, 51, 77, 78$ (the molecular ion at 78, loss of H at 77, C_4H_3^+ at 51, and C_2H_2^+ at 26) that correspond to genuine fragments of the structure rather than at arbitrary bins. The forward branch therefore produces molecule-specific spectra that are chemically meaningful, which is a prerequisite for using the forward model in retrieval.

4.2.2 Mass-Controlled Retrieval

The headline evaluation is spectrum-to-structure retrieval assessed against mass. Figure 4.2 reports, for held-out test molecules, the learned within-window MRR (mean reciprocal rank) next to the analytic random-chance MRR for two mass windows. Within ± 5 Da the learned retrieval reaches an MRR of 0.551 against a chance level of 0.503, a margin of +0.048; within ± 20 Da it reaches 0.219 against 0.206, a margin of +0.013. As these molecules were not observed during training or model selection, the positive margins suggest that the model is recovering structural information that generalises beyond the precursor mass.




Figure 4.2: Spectrum-to-structure retrieval on held-out test molecules, with precursor mass taken into account. For each mass window (± 5 and ± 20 Da around the query’s true mass), the learned encoder retrieval MRR is shown next to the analytic random-chance MRR for a pool of the same size. A bar standing above its chance reference indicates structural signal that mass alone cannot explain, since mass is approximately constant inside the window. This is the decisive test for structural signal: uncontrolled retrieval quality can be explained by mass alone, whereas here the model retains a positive margin over chance once mass is controlled for.

4.2.3 Why Raw Retrieval Scores Are Not the Right Axis

Mass dominates the uncontrolled metrics to such an extent that ranking candidates purely by their agreement with the query’s parent mass is already close to a solution in the full retrieval set. This is exactly what motivates the mass-controlled scheme. An oracle that ranks by the difference between the true mass and the mass implied by the query’s heaviest peak reaches an MRR of 0.92, and a spectrum-only mass baseline reaches 0.46. Against baselines this strong, the raw learned MRR over the full set, 0.11, does not clear this mass floor and so is not, on its own, evidence of learned chemistry. The full-set ablation figures below, which compare the full model against no-fragments, no-derivation-graph, and no-reranking variants are therefore reported as diagnostics of the pipeline’s behaviour rather than as the headline result.




Figure 4.3: Ablation of three pipeline components (fragments, derivation graph, and reranker) against the full model. Cosine similarity measures how similar the query is to the correct target, averaged over all queries; MRR is the mean reciprocal rank of the ground-truth target. Error bars show 95% bootstrap confidence intervals (5000 query resamples). Brackets indicate pairwise p-values from paired randomization tests computed separately for each metric. For *cosine similarity*, all variants are nearly identical (≈ 0.43) and no pairwise difference is statistically significant, so forward reconstruction is insensitive to the ablations. For *MRR*, removing reranking lowers the full-set score from 0.24 to 0.10 (Full vs. no-reranking, $p \approx 0.087$), whereas the fragment and derivation-graph ablations do not differ significantly from the full model. As discussed in the text, this full-set MRR is dominated by precursor mass, so the reranking effect shown here reflects mass-consistency filtering rather than learned spectral shape.

On the forward task, all evaluated variants achieve very similar cosine similarity values. The full model reaches a mean cosine similarity of 0.43, while the no-fragments and no-derivation-graph ablations achieve 0.43 and 0.43, respectively (Figure 4.3). These cosine values reflect genuine reconstructions; their near-equality across ablations suggests that the molecular encoder already captures most of the forward-predictive signal, with the fragment branch adding little on this particular metric.

For the full model, Recall@1 is 0.09, Recall@5 is 0.32, Recall@10 is 0.50, and Recall@20 is 0.86, with a mean reciprocal rank of 0.24 (Figure 4.4). Read on their own, these values resemble a system for narrowing down candidates, but for the reasons established above (Section 4.2.3)

they largely reflect agreement on precursor mass. Figure 4.2 isolates the genuinely learned component.




Figure 4.4: Recall@k for $k = 1, \dots, 20$ across the full retrieval model and three ablated variants. Shaded bands show 95% bootstrap confidence intervals over queries. The full model and the `no fragments` / `no derivation graph` variants cluster together across the full k range, while the `no reranking` variant trails the others at every k , with the largest gap at small k and gradual closure toward $k = 20$.

On the full set, removing reranking leaves the forward cosine similarity unchanged, as expected, but lowers retrieval quality: Recall@5 drops from 0.32 to 0.09, Recall@10 from 0.50 to 0.23, Recall@20 from 0.86 to 0.41, and MRR from 0.24 to 0.10. Taken at face value, this appears to be a significant advantage of forward-model rescoring. However, it is the mass-consistency filtering identified above (Section 4.2.3): the same improvement is achieved by retaining only the precursor-mass mask. Therefore, it does not provide evidence that the learned forward spectral shape drives retrieval. The fragment-related ablations remain inconclusive in terms of the full-set metrics. The differences between the full model and the no-fragments and no-derivation-graph variants are small. On query-level paired tests, these differences are not statistically significant. Therefore, the full-set view alone does not establish a benefit from the fragment branch. What the fragment branch does contribute unambiguously is the forward decoder’s molecule-specificity (Section 4.2.1), the property the mass-controlled retrieval result relies on.

These conclusions should be interpreted with caution. The held-out test set contains only 13 queries, so recall values change in coarse increments, and differences between variants often reflect only one or two queries. Model selection was performed on the validation split, so the

validation margins over chance are optimistic. The test set held out in Figure 4.2 is therefore more reliable. The numbers reported here come from a single training run and should be considered proof of concept rather than a benchmarked result.

Overall, the results demonstrate the effectiveness of the proposed bidirectional framework. Once the forward decoder is grounded in fragments, the retrieval of spectra from structures provides structural information beyond precursor masses for molecules that have not been used in training. Decisive evidence comes from the mass-controlled comparison rather than any raw retrieval score, since uncontrolled metrics — including the apparent benefit of reranking — are dominated by mass consistency. This work’s main contributions are twofold: a fragment-augmented bidirectional framework whose forward model is molecule-specific by design, and a mass-controlled evaluation that distinguishes between structural discrimination genuinely learned and mass filtering. Together, these provide a concrete basis for future work aimed at improving the informativeness and integration of fragment-level intermediate representations.

4.3 Challenges

The thesis is shaped by several interconnected challenges.

First, fragmentation must be represented in a way that is simultaneously chemically meaningful and useful for machine learning. *Second*, explicit fragment generation leads to a *combinatorial explosion* of candidate structures and limits the practical scale of the approach. *Third*, the absence of unique *ground-truth fragmentation graphs* makes fragment-level supervision inherently *ambiguous*. *Fourth*, bidirectional learning introduces *competing objectives*, since forward spectrum reconstruction and backward structure retrieval do not necessarily favor the same latent representation. Finally and *fifth*, the empirical contribution of fragment-aware intermediate representations remains difficult to establish conclusively, especially under *small-sample* conditions and in the presence of strong baseline signals from the molecular encoder alone.

Chapter 5

Future Work

A central limitation of the present framework is that it does not learn fragmentation itself. The model operates on *precomputed* derivation graphs. Candidate fragments and their parent–child relations are generated by the rule-based graph-transformation pipeline upstream, and are then encoded by the neural model. The learning problem thus begins only *after* a fragmentation space has been constructed. The network learns how to embed, combine, align, and score molecular, fragment, and spectral information, but not which fragmentation events to propose in the first place. This design is appropriate for a first proof of concept, since it keeps the fragment-level representation explicit and the learning problem tractable, but it also marks the clearest boundary of the current approach and thus the most important direction for future work.

5.1 From Fragment Encoding to Fragment Generation

The present implementation is best understood as a *fragment-aware encoder* rather than a full fragmentation simulator. Its fragment branch consumes fragment graphs together with their derivation structure encoding a graph whose nodes are fragment graphs and whose directed edges represent parent–child relations extracted from the precomputed derivation graph. The model therefore receives a structured fragmentation hypothesis as input instead of inferring it autonomously. The natural next step is to move from *encoding* fragmentation structures to *generating* them: rather than assuming the relevant fragment space has already been enumerated, the model would be trained to predict which bonds or substructures are most likely to fragment, in which order, and under which contextual conditions. This would shift the system from symbolic preprocessing with neural downstream learning toward a more integrated neuro-symbolic architecture, considerably closer to a true mechanistic simulator of fragmentation pathways.

5.2 Toward a Generate–Score Architecture

A key source of inspiration for this transition is ICEBERG [9], which separates the task into a *Generate* module that predicts likely fragmentation events and a *Score* module that assigns intensities to the resulting fragments. Crucially, the fragment graph is not treated as fixed input, but rather as a latent object that must be predicted itself. The generative stage is trained to reconstruct an estimated fragmentation hypergraph, while a second module predicts intensities and hydrogen rearrangements of the generated fragments. The present model already supplies the main ingredients for the *Score* side of such a system. It embeds molecular graphs, fragment-level structures, and spectra into a shared latent space; it predicts spectra from latent representations; and it uses bidirectional training signals such as contrastive alignment and reconstruction losses. What is missing is an upstream module that decides which fragments enter the fragment branch at all. The symbolic derivation graphs produced by the rule-based pipeline could be used to supervise such a generator — initially as a learnt pruning mechanism over the derivation graph and later as a more ambitious autoregressive model that predicts fragmentation events directly (Figure 5.1).

5.3 Advantages of Learning Fragment Generation

Learning fragmentation generation would offer several advantages over the present setup. First, it would reduce the dependence on exhaustive upstream preprocessing: instead of enumerating a full fragmentation space before learning begins, a learned generator could prioritize only the most informative or probable fragmentation events, thereby lowering the combinatorial burden and improving scalability. Second, it would make the architecture more *end-to-end*. Where the spectrum predictor is currently conditioned on a fragment representation that is treated as given, an integrated design could optimize fragment proposals, fragment embeddings, and spectrum scoring jointly, so the model learns not only how to use fragments but which fragment structures are most helpful for downstream spectrum reconstruction and inverse retrieval. Third, explicit fragmentation generation may improve interpretability. Because the fragment structures are currently externally generated, the model cannot explain why one fragmentation path was preferred over another, whereas a learned generative stage could provide atom- or bond-wise probabilities over fragmentation events and thereby give a more direct account of the model’s mechanistic assumptions.




Figure 5.1: A future extension of the present work: shifting from fragment-structure *encoding* to fragmentation-process *learning*, so that the fragmentation graph becomes a predicted object rather than a fixed input.

5.4 Scientific and Technical Challenges

This transition is non-trivial and raises several challenges. The first is supervision: spectral libraries generally provide molecule–spectrum pairs, but not ground-truth fragmentation graphs. As in ICEBERG, the derivation graphs produced by the rule-based MØD pipeline could serve as proxy labels for a learned generator, but such labels are not guaranteed to be chemically unique or physically exact, and this ambiguity would need to be addressed carefully during training and evaluation. The second is error propagation: once generation becomes learned, mistakes made by the generative stage may propagate into the spectral scoring stage. This argues for staged training, beam search, uncertainty estimation, or latent-variable formulations that can maintain several plausible fragmentation hypotheses rather than committing to a single one. The third is chemical realism: a learned generator should distinguish fragments that are structurally possible from those that are spectrometrically plausible, accounting for effects such as charge retention, neutral loss, and hydrogen rearrangements — again favoring a modular Generate–Score design over a single black-box network.

5.5 A Plausible Development Path

A realistic path forward would proceed in stages. The first stage would keep the current derivation graphs but train a neural module to rank or prune candidate fragmentation steps. The second stage would learn to expand fragmentation graphs autoregressively, conditioned on the precursor graph and the partial fragmentation history. The third stage would couple this generator tightly to the existing scoring model so that fragment generation, spectral prediction, and inverse retrieval are optimized jointly. Seen in this way, the present system is not an alternative to Generate–Score approaches but an intermediate step toward them: it demonstrates that fragment-level intermediate representations can strengthen the connection between molecules and spectra, and the next major advance is to let the model learn which fragmentation structures to construct rather than only how to encode those supplied by an external pipeline. This would bring the framework closer to a fully trainable, mechanistically informed, and end-to-end bidirectional model for structure–spectrum translation.

Bibliography

- [1] J. H. Gross, *Mass spectrometry, A textbook*, 2. printing. Berlin: Springer, 2006, 518 pp., ISBN: 9783540407393.
- [2] R. Cramer, “High-speed analysis of large sample sets - how can this key aspect of the omics be achieved?” *Molecular & Cellular Proteomics*, vol. 19, no. 11, pp. 1760–1766, Nov. 2020, ISSN: 1535-9476. DOI: [10.1074/mcp.p120.001997](https://doi.org/10.1074/mcp.p120.001997).
- [3] S. Stein, “Mass spectral reference libraries: An ever-expanding resource for chemical identification,” *Analytical Chemistry*, vol. 84, no. 17, pp. 7274–7282, Jul. 2012, ISSN: 1520-6882. DOI: [10.1021/ac301205z](https://doi.org/10.1021/ac301205z).
- [4] I. Blaženović, T. Kind, J. Ji, and O. Fiehn, “Software tools and approaches for compound identification of lc-ms/ms data in metabolomics,” *Metabolites*, vol. 8, no. 2, p. 31, May 2018, ISSN: 2218-1989. DOI: [10.3390/metabo8020031](https://doi.org/10.3390/metabo8020031).
- [5] A. Young, H. Röst, and B. Wang, “Tandem mass spectrum prediction for small molecules using graph transformers,” *Nature Machine Intelligence*, vol. 6, no. 4, pp. 404–416, Apr. 2024, ISSN: 2522-5839. DOI: [10.1038/s42256-024-00816-8](https://doi.org/10.1038/s42256-024-00816-8).
- [6] S. Goldman, J. Wohlwend, M. Stražar, G. Haroush, R. J. Xavier, and C. W. Coley, “Annotating metabolite mass spectra with domain-inspired chemical formula transformers,” *Nature Machine Intelligence*, vol. 5, no. 9, pp. 965–979, Aug. 2023, ISSN: 2522-5839. DOI: [10.1038/s42256-023-00708-3](https://doi.org/10.1038/s42256-023-00708-3).
- [7] Y. Nowatzky, F. F. Russo, J. Lisec, A. Kister, K. Reinert, T. Muth, and P. Benner, “Fiora: Local neighborhood-based prediction of compound mass spectra from single fragmentation events,” *Nature Communications*, vol. 16, no. 1, Mar. 2025, ISSN: 2041-1723. DOI: [10.1038/s41467-025-57422-4](https://doi.org/10.1038/s41467-025-57422-4).
- [8] F. F. Russo, Y. Nowatzky, C. Jaeger, M. K. Parr, P. Benner, T. Muth, and J. Lisec, “Machine learning methods for compound annotation in non-targeted mass spectrometry—a brief overview of fingerprinting, in silico fragmentation and de novo methods,”

- Rapid Communications in Mass Spectrometry*, vol. 38, no. 20, Aug. 2024, ISSN: 1097-0231. DOI: [10.1002/rcm.9876](https://doi.org/10.1002/rcm.9876).
- [9] S. Goldman, J. Li, and C. W. Coley, “Generating molecular fragmentation graphs with autoregressive neural networks,” *Analytical Chemistry*, vol. 96, no. 8, pp. 3419–3428, Feb. 2024, ISSN: 1520-6882. DOI: [10.1021/acs.analchem.3c04654](https://doi.org/10.1021/acs.analchem.3c04654).
 - [10] F. W. McLafferty and F. Tureček, *Interpretation of mass spectra*, eng, 4th ed. cau, Mill Valley, Calif, 1993, F.W. McLafferty, František Tureček., Includes bibliographical references (p. 321-338) and index., ISBN: 9780935702255.
 - [11] J. L. Andersen, C. Flamm, D. Merkle, and P. F. Stadler, *Generic strategies for chemical space exploration*, 2013. DOI: [10.48550/ARXIV.1302.4006](https://doi.org/10.48550/ARXIV.1302.4006).
 - [12] E. D. Hoffmann and V. Stroobant, *Mass spectrometry, Principles and applications*, 3. ed., reprinted. Chichester: Wiley, 2011, 489 pp., Includes bibliographical references and index, ISBN: 9780470033104.
 - [13] K. Scheubert, F. Hufsky, and S. Böcker, “Computational mass spectrometry for small molecules,” *Journal of Cheminformatics*, vol. 5, no. 1, Mar. 2013, ISSN: 1758-2946. DOI: [10.1186/1758-2946-5-12](https://doi.org/10.1186/1758-2946-5-12).
 - [14] S. R. Rambhad and H. Dave, “A review on applications of liquid chromatography coupled with mass spectrometry in impurity profiling of drug substances and drug products,” *SEPARATION SCIENCE PLUS*, vol. 8, no. 6, Jun. 2025, ISSN: 2573-1815. DOI: [10.1002/sscp.70073](https://doi.org/10.1002/sscp.70073).
 - [15] S. Böcker and F. Rasche, “Towards de novo identification of metabolites by analyzing tandem mass spectra,” *Bioinformatics*, vol. 24, no. 16, pp. i49–i55, Aug. 2008, ISSN: 1367-4803. DOI: [10.1093/bioinformatics/btn270](https://doi.org/10.1093/bioinformatics/btn270).
 - [16] A. Vaniya and O. Fiehn, “Using fragmentation trees and mass spectral trees for identifying unknown compounds in metabolomics,” *TrAC Trends in Analytical Chemistry*, vol. 69, pp. 52–61, Jun. 2015, ISSN: 0165-9936. DOI: [10.1016/j.trac.2015.04.002](https://doi.org/10.1016/j.trac.2015.04.002).
 - [17] K. Murray. “Mass spectrometry block schematic,” Wikimedia Commons. (May 2006), [Online]. Available: https://commons.wikimedia.org/wiki/File:Ms_block_schematic.gif (visited on 12/15/2025).
 - [18] J. Clark. “The mass spectrometer: How it works,” Chemguide. (2015), [Online]. Available: <https://www.chemguide.co.uk/analysis/masspec/masspec.GIF> (visited on 03/06/2026).

- [19] K. K. Murray, R. K. Boyd, M. N. Eberlin, G. J. Langley, L. Li, and Y. Naito, “Definitions of terms relating to mass spectrometry (iupac recommendations 2013),” *Pure and Applied Chemistry*, vol. 85, no. 7, pp. 1515–1609, Jun. 2013, ISSN: 0033-4545. DOI: [10.1351/pac-rec-06-04-06](https://doi.org/10.1351/pac-rec-06-04-06).
- [20] A. S. Moorthy, A. J. Kearsley, W. G. Mallard, W. E. Wallace, and S. E. Stein, “Inferring the nominal molecular mass of an analyte from its electron ionization mass spectrum,” *Analytical Chemistry*, vol. 95, no. 35, pp. 13 132–13 139, Aug. 2023, ISSN: 1520-6882. DOI: [10.1021/acs.analchem.3c01815](https://doi.org/10.1021/acs.analchem.3c01815).
- [21] A. L. Rockwood, M. M. Kushnir, and N. J. Clarke, “Mass spectrometry,” in *Principles and Applications of Clinical Mass Spectrometry*. Elsevier, 2018, pp. 33–65, ISBN: 9780128160633. DOI: [10.1016/b978-0-12-816063-3.00002-5](https://doi.org/10.1016/b978-0-12-816063-3.00002-5).
- [22] R. L. Stoddard, S. Collins, and J. S. McIndoe, *Mass spectrometry of organoaluminum derivatives*, Mar. 2016. DOI: [10.1002/9780470682531.pat0831](https://doi.org/10.1002/9780470682531.pat0831).
- [23] I. Ashyrmamatov, U. V. Ucak, and J. Lee, “Refining ei-ms library search results through atomic-level insights,” *Communications Chemistry*, vol. 8, no. 1, Nov. 2025, ISSN: 2399-3669. DOI: [10.1038/s42004-025-01706-9](https://doi.org/10.1038/s42004-025-01706-9).
- [24] C. Dass, *Fundamentals of contemporary mass spectrometry*, ser. Wiley-Interscience series on mass spectrometry. Hoboken, N.J.: Wiley-Interscience, 2007, 585 pp., Includes bibliographical references and index., ISBN: 9780470118481.
- [25] D. A. Vargas Medina, E. V. S. Maciel, N. G. Pereira dos Santos, and F. M. Lancas, “The overshadowed role of electron ionization–mass spectrometry in analytical biotechnology,” *Current Opinion in Biotechnology*, vol. 82, p. 102 965, Aug. 2023, ISSN: 0958-1669. DOI: [10.1016/j.copbio.2023.102965](https://doi.org/10.1016/j.copbio.2023.102965).
- [26] A. Amirav, A. B. Fialkov, A. Gordin, O. Elkabets, and K. J. Margolin Eren, “Cold electron ionization (ei) is not a supplementary ion source to standard ei. it is a highly superior replacement ion source,” *Journal of the American Society for Mass Spectrometry*, vol. 32, no. 11, pp. 2631–2635, Oct. 2021, ISSN: 1879-1123. DOI: [10.1021/jasms.1c00241](https://doi.org/10.1021/jasms.1c00241).
- [27] A. A. Ganeev, A. R. Gubal, S. V. Potapov, N. N. Agafonova, and V. M. Nemets, “Mass spectrometric methods for the direct elemental and isotopic analysis of solid material,” *Russian Chemical Reviews*, vol. 85, no. 4, pp. 427–444, Apr. 2016, ISSN: 1468-4837. DOI: [10.1070/rcr4504](https://doi.org/10.1070/rcr4504).

- [28] J. B. Fenn, M. Mann, C. K. Meng, S. F. Wong, and C. M. Whitehouse, "Electrospray ionization for mass spectrometry of large biomolecules," *Science*, vol. 246, no. 4926, pp. 64–71, Oct. 1989, ISSN: 1095-9203. DOI: [10.1126/science.2675315](https://doi.org/10.1126/science.2675315).
- [29] S. Banerjee and S. Mazumdar, "Electrospray ionization mass spectrometry: A technique to access the information beyond the molecular weight of the analyte," *International Journal of Analytical Chemistry*, vol. 2012, pp. 1–40, 2012, ISSN: 1687-8779. DOI: [10.1155/2012/282574](https://doi.org/10.1155/2012/282574).
- [30] M. Karas and F. Hillenkamp, "Laser desorption ionization of proteins with molecular masses exceeding 10,000 daltons," *Analytical Chemistry*, vol. 60, no. 20, pp. 2299–2301, Oct. 1988, ISSN: 1520-6882. DOI: [10.1021/ac00171a028](https://doi.org/10.1021/ac00171a028).
- [31] Y.-H. Lai and Y.-S. Wang, "Matrix-assisted laser desorption/ionization mass spectrometry: Mechanistic studies and methods for improving the structural identification of carbohydrates," *Mass Spectrometry*, vol. 6, no. 3, S0072–S0072, 2017, ISSN: 2187-137X. DOI: [10.5702/massspectrometry.s0072](https://doi.org/10.5702/massspectrometry.s0072).
- [32] Z. Bowen, S. An, and L. Zhang, "A mass spectrometer for measuring a high current ion beam with a big range of the charge-to-mass ratio," in *Proceedings of the 7th Int. Particle Accelerator Conf.*, vol. IPAC2016, Korea, 2016. DOI: [10.18429/JACOW-IPAC2016-WEPORO57](https://doi.org/10.18429/JACOW-IPAC2016-WEPORO57).
- [33] W. C. Wiley and I. H. McLaren, "Time-of-flight mass spectrometer with improved resolution," *Review of Scientific Instruments*, vol. 26, no. 12, pp. 1150–1157, Dec. 1955, ISSN: 1089-7623. DOI: [10.1063/1.1715212](https://doi.org/10.1063/1.1715212).
- [34] C. Li, S. Chu, S. Tan, X. Yin, Y. Jiang, X. Dai, X. Gong, X. Fang, and D. Tian, "Towards higher sensitivity of mass spectrometry: A perspective from the mass analyzers," *Frontiers in Chemistry*, vol. 9, Dec. 2021, ISSN: 2296-2646. DOI: [10.3389/fchem.2021.813359](https://doi.org/10.3389/fchem.2021.813359).
- [35] K.-H. Jang, S. Bae, J. Woo, S.-W. Kang, N.-S. Lee, and Y. U. Jeong, "Ion transmission probability through quadrupole mass filter depending on initial ion condition and rod misalignment," *AIP Advances*, vol. 12, no. 1, Jan. 2022, ISSN: 2158-3226. DOI: [10.1063/5.0076456](https://doi.org/10.1063/5.0076456).
- [36] A. Raffaelli and A. Saba, "Ion scanning or ion trapping: Why not both?" *Mass Spectrometry Reviews*, vol. 42, no. 4, pp. 1152–1173, Nov. 2021, ISSN: 1098-2787. DOI: [10.1002/mas.21746](https://doi.org/10.1002/mas.21746).
- [37] R. E. March, "Quadrupole ion traps," *Mass Spectrometry Reviews*, vol. 28, no. 6, pp. 961–989, Jun. 2009, ISSN: 1098-2787. DOI: [10.1002/mas.20250](https://doi.org/10.1002/mas.20250).

- [38] S. Bourcier and Y. Hoppilliard, "Use of diagnostic neutral losses for structural information on unknown aromatic metabolites: An experimental and theoretical study," *Rapid Communications in Mass Spectrometry*, vol. 23, no. 1, pp. 93–103, Dec. 2008, ISSN: 1097-0231. DOI: [10.1002/rcm.3852](https://doi.org/10.1002/rcm.3852).
- [39] M. Dantus, "Tracking molecular fragmentation in electron-ionization mass spectrometry with ultrafast time resolution," *Accounts of Chemical Research*, vol. 57, no. 6, pp. 845–854, Feb. 2024, ISSN: 1520-4898. DOI: [10.1021/acs.accounts.3c00713](https://doi.org/10.1021/acs.accounts.3c00713).
- [40] T. V. Harwood, D. G. C. Treen, M. Wang, W. de Jong, T. R. Northen, and B. P. Bowen, "Blink enables ultrafast tandem mass spectrometry cosine similarity scoring," *Scientific Reports*, vol. 13, no. 1, Aug. 2023, ISSN: 2045-2322. DOI: [10.1038/s41598-023-40496-9](https://doi.org/10.1038/s41598-023-40496-9).
- [41] M. B. Hamaneh, A. Y. Ogurtsov, O. I. Obolensky, and Y.-K. Yu, "Systematic assessment of deep learning-based predictors of fragmentation intensity profiles," *Journal of Proteome Research*, vol. 23, no. 6, pp. 1983–1999, May 2024, ISSN: 1535-3907. DOI: [10.1021/acs.jproteome.3c00857](https://doi.org/10.1021/acs.jproteome.3c00857).
- [42] J. Yasonik, "Multiobjective de novo drug design with recurrent neural networks and non-dominated sorting," *Journal of Cheminformatics*, vol. 12, no. 1, Feb. 2020, ISSN: 1758-2946. DOI: [10.1186/s13321-020-00419-6](https://doi.org/10.1186/s13321-020-00419-6).
- [43] J. N. Wei, D. Belanger, R. P. Adams, and D. Sculley, "Rapid prediction of electron-ionization mass spectrometry using neural networks," *ACS Central Science*, vol. 5, no. 4, pp. 700–708, Mar. 2019, ISSN: 2374-7951. DOI: [10.1021/acscentsci.9b00085](https://doi.org/10.1021/acscentsci.9b00085).
- [44] S. Bhattacharyya, R. Sahara, and K. Ohno, "A first-principles phase field method for quantitatively predicting multi-composition phase separation without thermodynamic empirical parameter," *Nature Communications*, vol. 10, no. 1, Aug. 2019, ISSN: 2041-1723. DOI: [10.1038/s41467-019-11248-z](https://doi.org/10.1038/s41467-019-11248-z).
- [45] V. I. Minkin, "Glossary of terms used in theoretical organic chemistry," *Pure and Applied Chemistry*, vol. 71, no. 10, pp. 1919–1981, Oct. 1999, ISSN: 0033-4545. DOI: [10.1351/pac199971101919](https://doi.org/10.1351/pac199971101919).
- [46] J.-H. Wang, W.-K. Choong, C.-T. Chen, and T.-Y. Sung, "Calibr improves spectral library search for spectrum-centric analysis of data independent acquisition proteomics," *Scientific Reports*, vol. 12, no. 1, Feb. 2022, ISSN: 2045-2322. DOI: [10.1038/s41598-022-06026-9](https://doi.org/10.1038/s41598-022-06026-9).
- [47] M. Ciach, B. Miasojedow, G. Skoraczyński, S. Majewski, M. Startek, D. Valkenborg, and A. Gambin, "Masserstein: Robust linear deconvolution by optimal transport," Jun. 2020. DOI: [10.1101/2020.06.02.129858](https://doi.org/10.1101/2020.06.02.129858).

- [48] F. Mémoli, “Gromov–wasserstein distances and the metric approach to object matching,” *Foundations of Computational Mathematics*, vol. 11, no. 4, pp. 417–487, Apr. 2011, ISSN: 1615-3383. DOI: [10.1007/s10208-011-9093-5](https://doi.org/10.1007/s10208-011-9093-5).
- [49] V. Mäkinen, “Peak alignment using restricted edit distances,” *Biomolecular Engineering*, vol. 24, no. 3, pp. 337–342, Sep. 2007, ISSN: 1389-0344. DOI: [10.1016/j.bioeng.2007.02.007](https://doi.org/10.1016/j.bioeng.2007.02.007).
- [50] Z. Wu, B. Ramsundar, E. N. Feinberg, J. Gomes, C. Geniesse, A. S. Pappu, K. Leswing, and V. Pande, “Moleculenet: A benchmark for molecular machine learning,” *Chemical Science*, vol. 9, no. 2, pp. 513–530, 2018, ISSN: 2041-6539. DOI: [10.1039/c7sc02664a](https://doi.org/10.1039/c7sc02664a).
- [51] S. Müller, C. Flamm, and P. F. Stadler, “What makes a reaction network “chemical”?” *Journal of Cheminformatics*, vol. 14, no. 1, Sep. 2022, ISSN: 1758-2946. DOI: [10.1186/s13321-022-00621-8](https://doi.org/10.1186/s13321-022-00621-8).
- [52] H.-C. Ehrlich and M. Rarey, “Systematic benchmark of substructure search in molecular graphs - from ullmann to vf2,” *Journal of Cheminformatics*, vol. 4, no. 1, Jul. 2012, ISSN: 1758-2946. DOI: [10.1186/1758-2946-4-13](https://doi.org/10.1186/1758-2946-4-13).
- [53] A. Wojtuch, T. Danel, S. Podlowska, and L. Maziarka, “Extended study on atomic featurization in graph neural networks for molecular property prediction,” *Journal of Cheminformatics*, vol. 15, no. 1, Sep. 2023, ISSN: 1758-2946. DOI: [10.1186/s13321-023-00751-7](https://doi.org/10.1186/s13321-023-00751-7).
- [54] S. Kearnes, K. McCloskey, M. Berndl, V. Pande, and P. Riley, “Molecular graph convolutions: Moving beyond fingerprints,” *Journal of Computer-Aided Molecular Design*, vol. 30, no. 8, pp. 595–608, Aug. 2016, ISSN: 1573-4951. DOI: [10.1007/s10822-016-9938-8](https://doi.org/10.1007/s10822-016-9938-8).
- [55] R. Echahed and M. Minas, Eds., *Graph Transformation, 9th International Conference, ICGT 2016, Held as Part of STAF 2016, Vienna, Austria, July 5–6, 2016, Proceedings*, ser. Lecture Notes in Computer Science. Springer International Publishing, 2016, vol. 9761, ISBN: 9783319405308. DOI: [10.1007/978-3-319-40530-8](https://doi.org/10.1007/978-3-319-40530-8).
- [56] A.-S. Abel, N. Lauber, J. L. Andersen, R. Fagerberg, D. Merkle, and C. Flamm, “Computational approaches in chemical space exploration for carbon fixation pathways,” *npj Systems Biology and Applications*, vol. 12, no. 1, Jan. 2026, ISSN: 2056-7189. DOI: [10.1038/s41540-025-00641-8](https://doi.org/10.1038/s41540-025-00641-8).
- [57] A. HABEL, J. MÜLLER, and D. PLUMP, “Double-pushout graph transformation revisited,” *Mathematical Structures in Computer Science*, vol. 11, no. 5, pp. 637–688, Sep. 2001, ISSN: 1469-8072. DOI: [10.1017/s0960129501003425](https://doi.org/10.1017/s0960129501003425).

- [58] C. Flamm, D. Merkle, and P. F. Stadler, “Computation in chemical graph rewriting networks,” *Journal of Physics: Complexity*, vol. 6, no. 1, p. 015 014, Mar. 2025, ISSN: 2632-072X. DOI: [10.1088/2632-072x/adb7b1](https://doi.org/10.1088/2632-072x/adb7b1).
- [59] J. L. Andersen, C. Flamm, D. Merkle, and P. F. Stadler, “An intermediate level of abstraction for computational systems chemistry,” *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, vol. 375, no. 2109, p. 20 160 354, Nov. 2017, ISSN: 1471-2962. DOI: [10.1098/rsta.2016.0354](https://doi.org/10.1098/rsta.2016.0354).
- [60] MØD Developers. “Derivation graph strategy (dgstrat) — mØd documentation.” Accessed 2026-03-24. (2022), [Online]. Available: <https://imada.sdu.dk/u/daniel/DM840-2022/Material/mod-doc/dgStrat/index.html> (visited on 03/24/2026).
- [61] W. E. Wallace and A. S. Moorthy, “Nist mass spectrometry data center standard reference libraries and software tools: Application to seized drug analysis,” *Journal of Forensic Sciences*, vol. 68, no. 5, pp. 1484–1493, May 2023, ISSN: 1556-4029. DOI: [10.1111/1556-4029.15284](https://doi.org/10.1111/1556-4029.15284).
- [62] M. Fey, J. Sunil, A. Nitta, R. Puri, M. Shah, B. Stojanovič, R. Bendias, A. Barghi, V. Kocijan, Z. Zhang, X. He, J. E. Lenssen, and J. Leskovec, *Pyg 2.0: Scalable learning on real world graphs*, 2025. DOI: [10.48550/ARXIV.2507.16991](https://doi.org/10.48550/ARXIV.2507.16991).
- [63] T. Böttcher, “An additive definition of molecular complexity,” *Journal of Chemical Information and Modeling*, vol. 56, no. 3, pp. 462–470, Feb. 2016, ISSN: 1549-960X. DOI: [10.1021/acs.jcim.5b00723](https://doi.org/10.1021/acs.jcim.5b00723).

Appendix A

Appendix

A.1 Hyperparameter Tables and Configurations

Table [A.1](#) lists the hyperparameters used to train the bidirectional model of Section [3.3.7](#). The tuned values (learning rate, weight decay, batch size, latent dimension, dropout, normalisation, the spectrum-loss and cycle-consistency weights, and the curriculum schedule) are the optimum selected by an Optuna hyperparameter search (20 trials); the remaining entries (network widths, epochs, spectrum binning) were fixed a priori.

Table A.1: Training hyperparameters. Tuned entries are the best trial of the Optuna sweep.

Setting	Symbol	Value
<i>Optimisation</i>		
Optimiser	—	Adam (one instance per phase)
Learning rate	η	7.1×10^{-5}
Weight decay	—	1.4×10^{-4}
Batch size	N	64
Epochs (Phase A + Phase B)	—	10000 + 10000
LR schedule / warmup / grad. clip	—	none
Checkpoint selection	—	final epoch
<i>Architecture</i>		
Shared latent dimension	d	256
GINEConv hidden width	d_h	128
GINEConv layers	—	3 (each $128 \rightarrow 128$; input proj. $18 \rightarrow 128$)
Derivation-graph GCN	—	2 GCNConv, $(d+1) \rightarrow 64 \rightarrow 256$, mean pool
Spectral MLP width	—	512
Dropout	—	0.43
Normalisation	—	BatchNorm
<i>Spectrum binning</i>		
m/z range	$[mz_{\min}, mz_{\max})$	$[1, 1000)$
Bin width	w_{bin}	1.0 Da
Number of bins	B	999
<i>Loss</i>		
Cosine weight	α_{cos}	1.59
Wasserstein weight	λ	1.17
Forward cycle weight	λ_{cyc}	0.77
Spectrum cycle weight	$\lambda_{\text{cyc-spec}}$	0.089
Spectrum–molecule align weight	λ_{align}	0.30
Contrastive temperature	τ	72 0.07
<i>Curriculum</i>		

A.2 Github Repository

The code for this thesis is available at <https://github.com/phireiser/mass-spec>

A.3 AI Use Statement

The author used the following AI tools during the writing of this thesis:

- **ChatGPT** (various versions) for language refinement and grammar checking throughout the document.
- **DeepL Write** for improving the style in the document.
- **Claude** for structuring and organizing.
- **GitHub Copilot** for code suggestions and auto-completion in the implementation of the codebase.